

OmniParse

AI-powered invoice extraction & verification

Legal Documents

Transparency, Terms, Privacy & Cookie Policy

VERSION

3.4 · Revised September 18, 2026

ISSUED BY

Simon Curda · Natural person · Czech Republic

CONTACT

damr58h@gmail.com

Table of Contents

This document consolidates the four legal instruments that govern your use of OmniParse. The AI Transparency Notice explains how the Service uses artificial intelligence. The Terms of Service set out the binding contract between you and Simon Curda. The Privacy Policy (including the embedded Data Processing Addendum) describes how personal data is handled. The Cookie Policy explains the storage used by the Service.

AI Transparency Notice

Last updated: September 18, 2026 · EU AI Act Article 50 disclosure

Terms of Service

Last updated: September 18, 2026 · Binding contract for use of the Service

Privacy Policy (incl. DPA)

Last updated: September 18, 2026 · GDPR compliance & Data Processing Addendum

Cookie Policy

Last updated: September 18, 2026 · Local storage & consent

Quick summary of v3.4 changes: This version aligns the legal documents with the actual production code behavior. **OpenRouter and Google Gemini (AI Studio free tier) are both disabled by default** in the production cascade (Mistral → Groq only); OpenRouter's free-tier models permit training on prompt content and have no DPA, while Google's AI Studio free tier is not covered by the Google Cloud DPA (only Vertex AI is). The Mistral integration now sends `usage_options = {"enable_training": false}` on every API call to disable training. The DPA has been expanded with §11.9 (Personal Data Breach Notification), §11.10 (Audit Rights), and §11.11 (Acceptance and Activation). ToS §10 Limitation of Liability has been restructured into a systematic framework (direct damages, indirect/consequential, aggregate cap, exceptions, mandatory consumer rights). ToS §7c adds an explicit 'No Payment Authorization' clause. The Privacy Policy includes a retention table distinguishing 30-day file-binary deletion from 90-day post-termination cleanup. Czech cookie-law citation corrected (Act No. 127/2005 Coll. § 89, not 480/2004 Sb.). Age of digital consent set to 15 (Czech Act No. 110/2019 Coll.). Separate EU withdrawal-acknowledgment checkbox and one-click withdrawal button implemented per Directive (EU) 2023/2673.

AI Transparency Notice

Last updated: September 18, 2026

EU AI Act Compliance (Regulation EU 2024/1689)

Per Article 50 of the EU AI Act (applicable since **2 August 2026**), deployers of AI systems must inform natural persons that they are interacting with an artificial intelligence system. This notice fulfills that obligation. Article 5 prohibitions on certain AI practices have been applicable since **2 February 2025**; OmniParse confirms that none of its processing falls within any Article 5 prohibited category (social scoring, manipulative AI, untargeted facial scraping, emotion recognition in workplaces/schools, etc.).

1. AI Disclosure

OmniParse uses artificial intelligence. Specifically:

- **Document parsing:** A vision-language model (VLM) analyzes uploaded documents (PDFs, images) to extract structured invoice data.
- **Chat assistant:** A large language model (LLM) generates responses to your questions about your invoice data.
- **Artifact generation:** The chat AI may generate structured data (tables, charts, summaries) based on your requests.

You are interacting with AI-generated output throughout the Service. All parsing results, chat responses, and generated artifacts are produced by AI models and should be treated accordingly.

2. System Classification

Based on the intended purposes and functionality of the Service, we have assessed OmniParse against the high-risk use cases listed in Annex III of the EU AI Act. **OmniParse is not a high-risk AI system** under Annex III, because invoice OCR/extraction and conversational analytics do not fall within any of the high-risk categories listed (which include biometric identification, critical infrastructure management, education/employment decisions, essential private/public services, law enforcement, migration/border control, and justice/democratic processes).

The Service is subject to applicable transparency requirements under Article 50 of the EU AI Act, in particular Article 50(1) (AI systems that interact with natural persons). The Service may also generate or manipulate content using AI. Where applicable, we will implement the transparency measures required by Article 50 for AI-generated or manipulated content (Article 50(4)). We do not assert that every output of the Service automatically falls within all Article 50(4) sub-categories; the applicability of specific transparency obligations is assessed based on the nature of the output and the context of its use.

As the provider of the Service under our own name and brand, we acknowledge our responsibilities under the EU AI Act. This assessment may change if the Service is used in high-risk contexts. We will update this notice accordingly and implement the required obligations if reclassification occurs.

General-Purpose AI (GPAI) classification: The AI models used by OmniParse (e.g., Gemini, Pixtral, Llama) are developed by upstream providers (Google, Mistral, Meta, etc.) and may qualify as GPAI models under Art. 3(63). OmniParse's AI Act role depends on the specific component and context: OmniParse may act as a

deployer of these models (within the meaning of Art. 3(4)) when it uses them under its own name for the invoice-extraction and chat-assistant purposes. For some components — particularly where OmniParse integrates the model into a broader service offered under its own brand and substantially shapes the system's intended purpose — OmniParse's role may also be that of a **downstream provider** of an AI system within the meaning of Art. 25. The exact classification of each component is fact-specific and should be confirmed by legal counsel. Transparency obligations applicable to GPAI providers under Art. 53 rest with the upstream model providers (Google, Mistral, Meta, etc.). OmniParse relies on the technical documentation and information provided by those upstream providers.

3. Technology Details

- **Type:** Document understanding and natural language processing
- **Components:** Vision-language model (VLM) for document parsing, Language model (LLM) for chat
- **Purpose:** Extract structured data from documents and provide conversational analysis
- **Data processed:** Invoices and related documents may contain personal data (e.g. names, email addresses, phone numbers, bank account details, signatures). Depending on how the Service is used, we may process personal data on behalf of users or their organizations.
- **AI providers:** We may use multiple third-party AI providers to process uploaded content. The provider used for a particular request may depend on availability, capacity, configuration, and other operational factors. A current list of subprocessors and applicable processing locations is available in our Privacy Policy.

4. Known Limitations

The AI may make mistakes. Specific risk areas include:

- **Inaccurate extraction:** Low-quality scans, unusual layouts, handwritten text, or non-English documents may produce errors.
- **Hallucinations in chat:** The AI may generate plausible but factually incorrect information about your data.
- **Artifact errors:** Generated tables, charts, and summaries inherit any extraction or reasoning errors.
- **Currency and formatting:** May misinterpret currency symbols, date formats, or number formatting across regions.
- **Calculation errors:** Calculations and statistics generated by the Service may contain errors, including errors caused by incorrect extraction, interpretation, aggregation, rounding, currency conversion or other processing. You are responsible for independently verifying results before relying on them.
- **Confidence indicators:** AI-generated confidence indicators are heuristic signals intended to help prioritize manual review. They are not statistical probabilities and do not constitute a guarantee of accuracy.
- **Provider availability:** AI providers may impose rate limits, discontinue models, or become unavailable. During periods of high demand, the Service may return "temporarily busy" messages.

Always verify AI-generated data against your original documents before using it for any purpose where accuracy matters (financial reporting, tax filing, legal proceedings).

5. No Professional Advice

The Service is provided for informational and administrative purposes only and does not constitute accounting, tax, legal, financial or other professional advice.

AI-generated calculations, classifications, summaries, recommendations and chat responses must not be relied upon as the sole basis for financial, accounting, tax, legal or business decisions. Always consult a qualified professional for advice specific to your situation.

Currency conversions, exchange rates, tax calculations and financial aggregations may be inaccurate or based on incomplete or outdated information.

6. Human Oversight

OmniParse provides the following human oversight mechanisms (consistent with AI Act Art. 14 best practices, even though the Service is not high-risk):

- **Confidence indicators:** Extractions include heuristic confidence indicators to help prioritize manual review. These are not statistical probabilities.
- **Manual review:** Users can review, edit, or delete any extracted data at any time.
- **No autonomous decisions:** The AI does not make decisions, categorize individuals, or take actions on its own. It only extracts and presents data for human review. Invoice approvals, payment authorizations, and vendor categorizations require explicit user action.
- **Feedback loop:** Users can re-upload documents, correct results, and delete inaccurate data.
- **Audit trail:** All user actions (edits, approvals, deletions) are logged in the user's account audit log for traceability.

7. Data Handling and AI Training

- **Processing by AI providers:** Documents and chat messages are sent to AI providers for inference. Data may be temporarily processed, cached, or retained by AI providers for security, abuse prevention, monitoring, billing, debugging, or other purposes specified in their applicable terms and data-processing agreements. Retention periods and processing locations vary by provider and service configuration.
- **Stored on our side:** We retain uploaded documents for up to 30 days in the active application environment. Backup copies, security logs and data processed by third-party providers may be retained for different periods where necessary for security, legal or operational purposes, as described in our Privacy Policy.
- **Model training — production configuration: OmniParse itself does not route customer content through AI API tiers that permit provider training on customer content** when the default configuration is in effect. Specifically: (a) we send `usage_options={"enable_training": false}` on every Mistral API call (effective on paid Mistral tier; on free tier, Mistral may ignore or reject this parameter) (b) OpenRouter (whose free-tier models typically permit training on prompt content) is **disabled by default**; (c) Groq's data-handling terms are governed by the Groq DPA and SCCs; (d) Google Gemini (AI Studio free tier) data-handling is governed by Google's AI Studio terms — for stronger guarantees, operators may migrate to Vertex AI (out of scope for this deployment). Operators who change this configuration must update the Privacy Policy and this AI Transparency Notice accordingly.
- **Important note on free-tier AI APIs:** Some AI providers may have different data handling terms for free/unpaid API tiers compared to paid tiers. We recommend reviewing the applicable provider terms for details. For production use involving personal data, paid API tiers may provide stronger data protection guarantees.
- **Metadata logging:** We log API request metadata (timestamps, success/failure, provider used) for operational purposes. We do not log the content of your documents to server logs.

8. Model and Provider Changes

AI models, providers and processing methods may be changed, updated, replaced or discontinued. Material changes to AI providers (e.g., switching the primary vision-extraction provider) will be announced in the Service or by email at least 7 days in advance, except where the change is urgently required for security or availability reasons. Changes in models may affect extraction results, classifications, calculations and chat responses. The same input document may produce different results when processed by different models or at different times.

The Service is provided on an "as is" and "as available" basis. We do not guarantee uninterrupted, error-free or continuously available operation of the Service or any AI provider. AI models and third-party providers may change, become unavailable, impose rate limits, discontinue models or modify their capabilities. Material changes will be notified per the above; immaterial changes (e.g., model version bumps within the same family) may happen without separate notice but will be reflected in the "Last updated" date of this Notice.

9. Your Rights as an AI System User

- Review original documents alongside extracted data
- Manually correct any AI-extracted fields
- Delete and re-upload documents to get new extractions
- Request information about the processing and factors relevant to an AI-generated result, where technically and legally available (contact us)
- Lodge a complaint about AI output quality (damr58h@gmail.com)
- You may stop using AI-powered processing and delete your data. Where supported, you may export your data for processing outside the Service.

10. International Transfers

AI processing is performed by Mistral AI (EU-based, Paris, France) and Groq Inc. (US — SCCs confirmed). **OpenRouter (US) and Google Gemini (US) are both disabled by default** and are not used unless the operator explicitly enables them after completing their own DPA/SCC review (OpenRouter) or AI Studio terms review (Google Gemini). Under GDPR Chapter V, transfers to the US require appropriate safeguards. We rely on Standard Contractual Clauses (SCCs) consistent with the Schrems II ruling for non-DPF-certified recipients, together with Transfer Impact Assessments (TIAs) where required. The EU-US Data Privacy Framework (DPF) adequacy decision (10 July 2023) covers DPF-certified US recipients — but note that Google's AI Studio free-tier endpoint is NOT covered by the Google Cloud DPA (which applies to the Vertex AI path only this is why Google Gemini is disabled by default in our configuration).

***SCC Status by provider (as of September 18, 2026): Mistral AI** — EU-based (Paris, France). Transfers are expected to remain within the EEA, subject to Mistral's applicable terms. No SCC required. Training opt-out (usage_options.enable_training=false) sent on every request when MISTRAL_DISABLE_TRAINING=true (default). **Groq Inc.** — SCCs confirmed in effect (DPA dated October 15, 2025; EU SCC Module 2 self-executing upon acceptance of Groq Services Agreement). DPA accessible at console.groq.com (<https://console.groq.com>). Governing law: Ireland. Competent authority: Irish DPC. 72-hour breach notification. **Google LLC (Gemini)** — **DISABLED BY DEFAULT**. Uses AI Studio free-tier endpoint (generativelanguage.googleapis.com). The Google Cloud DPA (cloud.google.com/terms/data-processing-addendum) applies to the Vertex AI path, NOT the AI Studio free tier. Google LLC is DPF-certified, but the AI Studio free-tier data-handling terms are weaker than Vertex AI.*

Operator must set `ENABLE_GOOGLE_GEMINI=true` after reviewing AI Studio terms. For EU personal data, migrate to Vertex AI. **Vercel Inc.** — DPF-certified. DPA available at vercel.com/legal/dpa (<https://vercel.com/legal/dpa>). **Stripe Inc.** — DPF-certified. DPA available at stripe.com/legal/dpa (<https://stripe.com/legal/dpa>). **OpenRouter** — **DISABLED BY DEFAULT**. No personal data is transferred to OpenRouter unless the operator explicitly sets `ENABLE_OPENROUTER=true`. Before enabling, the operator must complete a DPA with OpenRouter, verify SCCs or DPF certification, switch to paid-tier models (to disable training), and update this notice. **Summary:** AI processing via Mistral (EU) and Groq (US, SCCs confirmed) has appropriate safeguards in place. OpenRouter and Google Gemini are disabled by default and are not used unless the operator explicitly enables them. Users should note that documents may contain personal data of multiple data subjects — the user is responsible for ensuring a valid legal basis for processing and transferring such data.

Your primary database (Supabase) is hosted in Ireland (EU) — no transfer outside the EU for the primary data store.

11a. AI Literacy (AI Act Art. 4)

Pursuant to Article 4 of the EU AI Act (applicable since 2 February 2025), providers and deployers of AI systems shall take measures to ensure, to their best extent, a sufficient level of AI literacy in their staff and persons dealing with the operation and use of AI systems on their behalf. This Notice forms part of OmniParse's measures supporting AI literacy among persons involved in the operation and use of the Service. It is not, by itself, a complete fulfillment of the Article 4 obligation — that obligation also encompasses internal staff training, operational documentation, and ongoing assessment of AI literacy needs. Internal staff training on AI capabilities, limitations, and risks is documented in our internal records and updated at least annually, or when significant changes to AI features occur.

11b. Logging and Monitoring

Although OmniParse is not classified as a high-risk AI system (and therefore Article 12 logging requirements do not strictly apply), we maintain the following logging practices as best practice:

- **Application-level audit log:** User actions (invoice create/edit/approve/delete, chat send) are logged in the user's account and retained for the lifetime of the account.
- **Server-side operational logs:** Vercel function logs retained for 30 days for debugging and incident investigation.
- **AI request metadata:** Timestamps, provider used, success/failure status, and request duration are logged for operational monitoring. The **content** of documents and chat messages is not logged to server logs.
- **Incident logs:** Security incidents involving AI systems are documented and retained for at least 2 years for regulatory audit purposes.

12. Contact

- **AI ethics concerns:** damr58h@gmail.com (subject line "AI Ethics")
- **Data protection / GDPR:** damr58h@gmail.com (subject line "DSR" or "DPA")
- **Legal:** damr58h@gmail.com
- **Czech DPA (ÚOOÚ):** [uoou.cz](https://www.uoou.cz) (<https://www.uoou.cz>)

The Service is intended to comply with applicable laws in the jurisdictions in which it is offered. Users remain responsible for ensuring that their use of the Service complies with applicable local requirements.

Terms of Service

Last updated: September 18, 2026

1. Acceptance of Terms

By accessing or using OmniParse ("the Service"), you agree to be bound by these Terms of Service ("Terms"). If you do not agree with any part of these Terms, you must not use the Service. These Terms apply to all users of the Service.

When you register an account, you are required to explicitly accept these Terms by checking the "I agree to the Terms of Service and Privacy Policy" checkbox. The timestamp of your acceptance is recorded in your account record. Continued use of the Service after any changes to these Terms constitutes acceptance of the updated Terms.

2. Description of Service

OmniParse is an AI-powered document parsing application that extracts structured data from invoices and similar business documents. The Service uses vision-language models to analyze uploaded documents and return extracted data in JSON format.

Current capabilities:

- Document upload (PDF, JPEG, PNG, WebP) up to 20MB per file
- AI-powered field extraction with per-field confidence scores
- AI chat assistant with interactive artifact generation (tables, charts, summaries)
- Data export (CSV and JSON) from the invoices table
- Visual analytics dashboard with charts
- User authentication with email, password, and JWT tokens
- GDPR-compliant data export and account deletion

3. Eligibility

You must be at least **15 years old** to use this Service. This age corresponds to the age of digital consent in the Czech Republic under GDPR Art. 8 (as implemented by Czech Act No. 110/2019 Coll.). In other EU/EEA Member States where the age of digital consent is higher (e.g., 16 in Germany), you must meet that higher age requirement. The Service is not directed to children, and we do not knowingly collect personal data from children below the applicable age of digital consent.

If you are under 18 but above the applicable age of digital consent, you confirm that you have obtained parental or guardian consent to use this Service and to enter into these Terms.

If you are registering on behalf of a business entity, you confirm that you have the authority to bind that entity to these Terms.

4. User Responsibilities and Warranties

You represent and warrant that you have all rights, permissions and lawful authority necessary to upload, process and transmit documents, emails and other content through the Service.

- You are solely responsible for the documents you upload to the Service.
- You must have the legal right to process any documents you submit (ownership, authorization, or legitimate interest).
- You must not upload documents containing personal data of EU data subjects unless you have a legal basis to do so under GDPR.
- Unless explicitly supported by the Service, users must not upload documents containing special categories of personal data (such as health information, biometric data, information revealing religious or political beliefs, or other highly sensitive information).
- When using email import, the Service may process email content, headers, attachments and other information necessary to identify and process invoices. Users should not forward emails containing information that they are not authorized to disclose to the Service.
- You must not use the Service for any unlawful purpose, including fraud, money laundering, or tax evasion.
- You must not attempt to reverse-engineer, decompile, or disrupt the Service.
- You must not use automated tools to overwhelm the Service (no scraping, no DDoS, no rate-limit evasion).
- You must not attempt to extract other users' data through IDOR attacks or similar exploits.
- You acknowledge that AI-extracted data may contain errors and you are responsible for verifying results before use.
- You are responsible for keeping your account password confidential.
- You are responsible for maintaining appropriate backups of documents and data that are important to you. The Service should not be considered the sole repository or archival system for financial records.

5. Intellectual Property and License

Your content: Content you upload (documents, chat messages, extracted data) remains your property. OmniParse does not claim ownership of your documents or extracted data. By uploading content, you grant OmniParse a worldwide, non-exclusive, royalty-free, revocable license to process, transmit, store, and display your content solely for the purpose of operating and providing the Service to you, including the right to transmit content to third-party AI inference providers as described in the Privacy Policy. This license ends when the relevant content is deleted from the Service, except to the extent continued processing or retention is required for legitimate operational, security, legal, or backup purposes described in the Privacy Policy (including backup copies that may persist for a limited period after deletion from active storage).

Service IP: The Service itself, including its source code, design, UI, documentation, brand, and underlying infrastructure, is the property of Simon Curda. You may not copy, modify, distribute, or create derivative works from the Service itself without written permission. The names "OmniParse" and associated logos are trademarks of Simon Curda.

Feedback: If you provide feedback, suggestions, or ideas about the Service, you grant OmniParse a perpetual, irrevocable, royalty-free license to use and incorporate that feedback into the Service without obligation or compensation to you.

6. Data Processing

Uploaded documents are retained in the active application environment for up to 30 days for preview purposes and are then deleted from active storage, subject to applicable backup, security, legal and provider retention periods described in our **Privacy Policy** (/privacy-policy) (Privacy Policy §4 retention table). Extracted data is stored in your account database and is accessible from your dashboard. You can delete your data at any time. See our **Privacy Policy** (/privacy-policy) for full details on data handling, including the distinction between active-storage deletion (30 days) and post-termination cleanup (90 days).

7. AI Disclaimer

THE SERVICE USES ARTIFICIAL INTELLIGENCE FOR DOCUMENT ANALYSIS, DATA EXTRACTION, AND CONVERSATIONAL CHAT RESPONSES. AI OUTPUT MAY CONTAIN ERRORS, INACCURACIES, OR HALLUCINATIONS.

You must always verify AI-extracted data against the original documents before using it for financial reporting, tax filing, legal proceedings, or any context where accuracy is critical. OmniParse provides confidence scores as a guide, but these are estimates and not guarantees of accuracy.

The AI Chat assistant is a conversational tool, not a source of truth. Responses to your questions about your invoice data are generated by AI models and may be inaccurate, incomplete, or out-of-context. Always cross-reference chat responses against the underlying invoice records and source documents before making financial decisions, payments, vendor communications, or tax filings based on them.

AI providers: OmniParse uses third-party AI providers for vision extraction and text generation. The primary provider is **Mistral AI** (Paris, France — EU-based; transfers are expected to remain within the EEA, subject to Mistral's applicable terms). **Groq Inc.** (US-based) has confirmed Standard Contractual Clauses (SCCs) in effect (EU SCC Module 2, self-executing under Groq's DPA dated October 15, 2025). **OpenRouter** (US-based) and **Google Gemini** (US-based, AI Studio free tier) are **disabled by default** and are not used unless the operator explicitly enables them after completing their own DPA/SCC review (OpenRouter) or AI Studio terms review (Google Gemini). These providers may experience rate limits, outages, or model deprecations that affect the availability and quality of AI features. We do not control and are not responsible for the behavior of these upstream models beyond our integration layer.

See our **AI Transparency Notice** (/ai-act-notice) for full details on AI limitations and your rights.

7a. No Professional Advice

The Service is provided for informational and administrative purposes only and does not constitute accounting, tax, legal, financial or other professional advice. AI-generated calculations, classifications, summaries, recommendations and chat responses must not be relied upon as the sole basis for financial, accounting, tax, legal or business decisions.

Currency conversions, exchange rates, tax calculations and financial aggregations may be inaccurate or based on incomplete or outdated information. Always consult a qualified professional for advice specific to your situation.

7b. Calculation and Data Accuracy Disclaimer

Calculations and statistics generated by the Service may contain errors, including errors caused by incorrect extraction, interpretation, aggregation, rounding, currency conversion or other processing. You are responsible for independently verifying all results before relying on them for any purpose.

AI models, providers and processing methods may be changed, updated, replaced or discontinued. Material changes to AI providers (e.g., switching the primary vision-extraction provider) will be announced in the Service or by email at least 7 days in advance, except where the change is urgently required for security or availability reasons. Changes in models may affect extraction results, classifications, calculations and chat responses. The same input document may produce different results when processed by different models or at different times.

7c. No Payment Authorization

OmniParse is an invoice intelligence platform, not a payment system. Unless explicitly enabled by you through a supported integration (which OmniParse does not currently offer), OmniParse does not initiate, authorize, execute, or otherwise cause any payment, bank transfer, wire transfer, ACH transaction, credit-card charge, or other financial transaction. The Service extracts, organizes, and analyzes invoice data — it does not pay invoices.

Any payment you make to a vendor, supplier, or other recipient — even if OmniParse extracted the vendor's bank details, suggested a payment amount, or displayed a "Pay" button that links to an external site — is solely your responsibility. You must independently verify the payee, amount, bank account, and payment method before initiating any payment. To the maximum extent permitted by applicable law, OmniParse is not a party to any payment transaction and bears no liability for payments made, missed, or incorrectly executed based on data extracted or suggested by the Service. Nothing in this section excludes or limits liability that cannot lawfully be excluded or limited under applicable law, including liability for intentional damage, gross negligence, and mandatory consumer protection rights (see §10.4 and §10.5).

If OmniParse adds payment-integration features in the future (e.g., bank integrations, accounting-system integrations, payment links), such features will be governed by separate terms specific to those integrations and will require your explicit opt-in. This section applies to the current Service, which does not include payment-execution capabilities.

8. Service Availability and Limitations (Including IMAP Email Capture)

The Service is provided on a best-effort basis. OmniParse uses automated processes (including IMAP email scanning, AI vision-language models, classification heuristics, and external AI providers) to capture and extract invoice data. None of these processes are guaranteed to be complete, accurate, or timely. We explicitly disclaim any warranty that the Service will capture every invoice, classify it correctly, or extract all data without error.

The IMAP email auto-capture feature is an auxiliary convenience, not a substitute for proper invoice management. Specifically, the Service may fail to capture invoices for reasons including but not limited to:

- AI provider rate limits, outages, model deprecations, or temporary unavailability
- Classifier false negatives (the AI incorrectly determines an email is not an invoice)
- Email attachment format issues (corrupted PDFs, unsupported image formats, encrypted archives)
- IMAP server-side issues (mailbox sync delays, throttling by Gmail/Outlook/etc.)

- Vercel function timeout (60-second limit per scan call)
- Network connectivity issues between Vercel, the email provider, and AI providers
- Pending Review queue cap (100 items per inbox — scan pauses when reached)
- Sender blocklist (user or system blocks a legitimate invoice sender)
- Emails older than the scan cursor (UID-based resume may skip emails if cursor is corrupted)
- Bugs, configuration errors, or schema migrations that affect data integrity

You must not rely solely on the Service for time-critical or financially significant invoice management.

You remain responsible for:

- Periodically checking that all expected invoices have been captured
- Manually uploading invoices that the Service failed to capture
- Verifying due dates and payment deadlines against source documents
- Maintaining backup processes for invoice tracking independent of OmniParse
- Reviewing the Pending Review queue regularly (at minimum weekly, or before any payment run)

To the maximum extent permitted by applicable law, OmniParse shall not be liable for losses resulting from missed, delayed, incorrectly classified, or incorrectly extracted invoices where such losses arise from limitations, errors, or failures of the Service. This includes late payment fees, vendor relationship damage, tax filing errors, late filing penalties, or any other damages arising from the Service's failure to capture or correctly process an invoice, including scenarios where the user has stopped checking their email inbox manually and relies exclusively on the Service's auto-capture feature. Nothing in this section excludes or limits liability that cannot lawfully be excluded or limited under applicable law, including liability for intentional damage, gross negligence, and mandatory consumer protection rights (see §10.4 and §10.5).

For paid users only: Where OmniParse is found directly at fault for damages (e.g., a confirmed bug in our code, not a third-party AI provider issue), liability is capped at the amount you paid to OmniParse in the 12 months preceding the claim. For Free tier users, liability for slight negligence is excluded to the extent permitted by applicable law; liability for intentional damage and gross negligence remains unaffected in all cases. Nothing in this section excludes or limits liability that cannot lawfully be excluded or limited under applicable law, including mandatory consumer-protection rights under Czech Act No. 634/1992 Coll. and EU Directive 93/13/EEC.

Notification of issues: If you become aware that the Service has failed to capture an invoice or has extracted incorrect data, please notify us at damr58h@gmail.com as soon as practicable. We will investigate and, where possible, restore or correct the affected data. Prompt notification helps us serve you better; this does not shorten any statutory limitation period under applicable law.

9. Disclaimer of Warranties

THE SERVICE IS PROVIDED "AS IS" AND "AS AVAILABLE" WITHOUT WARRANTIES OF ANY KIND, EITHER EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO IMPLIED WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE, AND NON-INFRINGEMENT.

We do not warrant that the Service will be uninterrupted, error-free, or free of harmful components. We do not warrant that AI extraction results will be accurate or complete. AI providers (Mistral, Groq) may experience rate limits, model deprecations, or outages that affect availability of AI features. Mistral is EU-based (Paris, France) Groq is US-based. OpenRouter and Google Gemini are **disabled by default** and are not used unless explicitly enabled by the operator after completing their own DPA/SCC review (OpenRouter) or AI Studio

terms review (Google Gemini).

10. Limitation of Liability

This Section 10 sets out the limitation of liability framework between you and Simon Curda (the operator of OmniParse). It is structured as: (a) direct damages, (b) indirect/consequential damages, (c) aggregate liability cap, (d) exceptions, (e) mandatory statutory rights.

10.1 Direct Damages

To the maximum extent permitted by applicable law, Simon Curda's liability for direct damages arising from your use of the Service is capped at the amount you paid to OmniParse in the 12 months preceding the claim. For Free tier users (who paid nothing), liability for direct damages is capped at zero to the extent permitted by applicable law.

10.2 Indirect and Consequential Damages

To the maximum extent permitted by applicable law, Simon Curda shall not be liable for any indirect, incidental, special, consequential, or punitive damages, including but not limited to loss of profits, loss of data, loss of business opportunities, reputational damage, or costs of substitute services. This includes losses arising from reliance on AI-generated output that you failed to independently verify against source documents.

10.3 Aggregate Liability Cap

The aggregate liability of Simon Curda for all claims arising out of or relating to these Terms or the Service (whether in contract, tort, or otherwise) shall not exceed the total amount you paid to OmniParse in the 12 months preceding the claim. For Free tier users, the aggregate liability cap is zero to the extent permitted by applicable law.

10.4 Exceptions (Liability Not Limited)

The limitations in §10.1, §10.2, and §10.3 do **not** apply to liability that cannot lawfully be excluded or limited under applicable law, including:

- Liability for intentional damage or intentional misconduct;
- Liability for gross negligence under applicable mandatory provisions of Czech law;
- Liability for personal injury or death caused by negligence;
- Liability arising from breach of statutory duties under mandatory consumer-protection law;
- Liability for defects in goods or services to the extent such liability cannot be contractually excluded under applicable mandatory law;
- Liability under product liability laws;
- Any other liability that cannot be excluded or limited under applicable mandatory law.

10.5 Mandatory Consumer Rights (EU / Czech)

EU consumer note: If you are a consumer within the meaning of Czech Act No. 89/2012 Coll. (Civil Code) § 419, nothing in these Terms limits any statutory consumer rights you may have under EU or Czech law, including under:

- Czech Act No. 89/2012 Coll. (Civil Code) — contractual liability, defects, consumer protection;
- Czech Act No. 634/1992 Coll. (Consumer Protection Act) — out-of-court dispute resolution;
- EU Directive 93/13/EEC (Unfair Contract Terms) — unfair terms are not binding on consumers;

- EU Directive 2011/83/EU (Consumer Rights) — right of withdrawal, pre-contractual information;
- GDPR (EU Regulation 2016/679) — data subject rights, liability under Art. 82.

Where mandatory consumer-protection law applies, it overrides conflicting provisions of these Terms to the extent of the conflict. If any provision of §10.1–10.3 is found unenforceable against a consumer, that provision shall be limited or eliminated to the minimum extent necessary to comply with applicable law, and the remaining provisions shall remain in full force and effect.

10.6 Basis of Limitation

The limitations in this Section 10 reflect the allocation of risk between the parties and form an essential basis of the bargain. The Service is offered at a price that reflects these limitations, and you acknowledge that the price would be substantially higher absent these limitations. Operators who require higher liability caps (e.g., Enterprise customers with negotiated contracts) may contact damr58h@gmail.com to discuss custom terms.

11. Indemnification

To the extent permitted by applicable law, and except where prohibited by mandatory consumer protection law (including EU Directive 93/13/EEC on unfair contract terms and Czech Act No. 634/1992 Coll.), you agree to indemnify and hold harmless Simon Curda from and against any claims, damages, losses, or expenses (including reasonable attorneys' fees) arising from (a) your use of the Service in violation of these Terms or applicable law, (b) your infringement of any third-party rights (including uploading documents you do not have the right to process), or (c) your violation of applicable data-protection law in respect of personal data you upload.

Consumer carve-out: If you are a consumer within the meaning of Czech Act No. 89/2012 Coll. (Civil Code) § 419, this indemnification clause applies only to the extent that the relevant claim arises from your intentional breach or gross negligence, and in no event shall a consumer be required to indemnify Simon Curda for damages caused by Simon Curda's own act or omission.

12. Subscription Plans and Billing

Paid plans (Pro, Plus, Business, Enterprise) are billed monthly or annually via Stripe. You can cancel your subscription at any time from your Stripe customer portal. Cancellation takes effect at the end of the current billing period — you retain access to paid features until then, and no further charges are made.

When you delete your account, any active Stripe subscription is automatically cancelled to prevent ongoing charges.

Refund policy:

- **Within 14 days of payment, no AI usage:** Full refund on request (see §14 Right of Withdrawal). Contact damr58h@gmail.com.
- **Within 14 days of payment, AI features used:** Right of withdrawal is forfeited under §1837(j) of the Civil Code (digital content with express consent to start). No refund, but you retain access until the end of the billing period.
- **After 14 days, monthly plan:** No refund of the current month's charge; subscription remains active until end of billing period. No further charges after cancellation.
- **After 14 days, annual plan:** Pro-rata refund of unused full months remaining on the plan, less a reasonable administrative fee of EUR 10. Contact damr58h@gmail.com to request a pro-rata refund.

- **Service-impacting outage:** If OmniParse experiences a verified outage exceeding 24 consecutive hours that prevents you from using paid features, contact us for a service-credit or pro-rata refund for the affected period.
- **Free tier:** No payment is made, so no refund is applicable. Account deletion is available at any time.

Refunds are processed back to the original payment method via Stripe within 14 days of approval, as required by Art. 13(1) of Directive 2011/83/EU.

13. Modifications to Terms

We may update these Terms from time to time. **Material changes will be notified to you via email (to your account email address) and via an in-Service notice at least 30 days before the change takes effect.** For changes that are not material (e.g., corrections of typos, clarifications, or changes that do not affect your rights or obligations), we may update these Terms without prior notice, with effect from the "Last updated" date.

For consumer users: Pursuant to § 1752 of the Civil Code (Act No. 89/2012 Coll.), material changes to these Terms that would worsen your position as a consumer take effect only if you expressly agree to them. Your continued use of the Service after the 30-day notice period constitutes such express agreement unless you terminate your account before the change takes effect (in which case the previous version of the Terms continues to apply for any open billing period).

The "Last updated" date at the top of this page indicates when these Terms were last revised. We will retain prior versions of these Terms and provide them on request.

14. Right of Withdrawal (EU Consumers)

Under Czech Civil Code (Act No. 89/2012 Coll.) §§ 1829 and 1837, and Directive 2011/83/EU on consumer rights, consumers concluding distance contracts (contracts concluded online) have a **14-day right of withdrawal**.

However, pursuant to § 1837(j) of the Civil Code and Art. 16(m) of Directive 2011/83/EU, the right of withdrawal does **NOT** apply to:

- Digital content or services where the consumer has expressly requested that the service begin during the withdrawal period AND has acknowledged that by doing so, they lose the right to withdraw;
- Services that have been fully performed with the consumer's prior express consent and acknowledgment.

For this waiver to be valid under § 1837(j) and Art. 16(m) CRD, the following express acknowledgment is required at signup (separate from the Terms acceptance checkbox):

"I acknowledge that OmniParse is a digital service that begins immediately upon my first use of AI features (upload, scan, or chat). I expressly consent to OmniParse beginning performance during the 14-day withdrawal period, and I acknowledge that I will lose my right of withdrawal once the service has started, pursuant to § 1837(j) of the Czech Civil Code and Art. 16(m) of Directive 2011/83/EU."

By creating an OmniParse account and checking this separate acknowledgment box, you expressly agree that:

- The Service begins immediately upon your first use of AI features (upload, scan, or chat)
- You acknowledge that you are losing the 14-day right of withdrawal pursuant to § 1837(j) of the Civil Code by requesting the Service to begin during the withdrawal period;
- This waiver applies only to the AI processing features. You may still delete your account at any time (Settings → Delete Account) without obligation.

One-click withdrawal button: Pursuant to Art. 4(2) of Directive (EU) 2023/2673 (effective 19 June 2026), if you have not yet used any AI features, a one-click "Withdraw" button is available in your account Settings that allows you to exercise your right of withdrawal without further formalities. Withdrawal must be confirmed by us within 24 hours.

For Free tier users: Since no payment is made, no withdrawal refund is applicable. Account deletion is available at any time.

For paid subscribers: If you cancel your subscription within 14 days of payment AND have not used any AI features in that period, you are entitled to a full refund (processed within 14 days via Stripe). Contact damr58h@gmail.com to request a refund. If you have used AI features during the 14-day period, the right of withdrawal is forfeited per § 1837(j).

15. Consumer Information and Dispute Resolution

Pursuant to § 1753 of the Civil Code (Act No. 89/2012 Coll.), the following information is provided to consumers before concluding a distance contract:

- **Provider:** Simon Curda (natural person), Czech Republic
- **Contact:** damr58h@gmail.com
- **Service description:** AI-powered invoice extraction, validation, and analytics platform
- **Total price (incl. taxes):** Free tier available (no charge). Paid plans from USD 49/month to USD 499/month as described on the pricing page; VAT is reverse-charged for EU B2B customers with a valid VAT ID and included for EU consumers at the applicable Czech VAT rate.
- **Additional fees:** None beyond the subscription price. No delivery costs (digital service delivered online).
- **Contract duration:** Monthly or annual subscription, cancelable at any time (cancellation takes effect at end of billing period).
- **Performance timeframe:** The Service begins immediately upon account creation and first use of AI features.
- **Payment method:** Credit/debit card via Stripe. Invoicing available for Enterprise plans on request.
- **Right of withdrawal:** See § 14 above (limited for digital content under § 1837(j) Civil Code).
- **Applicable law:** Czech law (see § 16 below)
- **Court of jurisdiction:** For consumer disputes, the court of the consumer's domicile; for B2B disputes, the competent courts of the Czech Republic.
- **Language:** These Terms are available in English. Communication with the provider may be conducted in English or Czech.
- **After-sales obligations:** None beyond the warranties and obligations set out in these Terms.
- **Professional liability insurance:** Not maintained (the operator is a natural person acting in a non-commercial capacity claims are subject to the liability caps in § 10).
- **Code of conduct:** The provider has not subscribed to any code of conduct under § 1753(1)(i) Civil Code.

Out-of-Court Dispute Resolution (ADR)

Czech consumers may seek out-of-court dispute resolution through the following entity:

- **Česká obchodní inspekce (ČOI)** — Czech Trade Inspection Authority
- Address: Štěpánská 567, 120 00 Praha 2, Czech Republic

- Website: **coi.gov.cz** (<https://coi.gov.cz>)
- ADR information: **coi.gov.cz/informace-o-adr** (<https://coi.gov.cz/informace-o-adr>)

A proposal for out-of-court dispute resolution must be submitted within 1 year from the date the consumer first exercised their right in the matter. The procedure is governed by Act No. 634/1992 Coll. (Consumer Protection Act) §§ 20b–20e.

EU consumers may also seek alternative dispute resolution through the European Commission's consumer redress resources at **consumer-redress.ec.europa.eu** (<https://consumer-redress.ec.europa.eu>).

16. Governing Law and Dispute Resolution

These Terms are governed by the laws of the Czech Republic, without regard to conflict of law principles. EU consumer protection laws apply additionally for users classified as consumers under EU law.

For disputes arising from these Terms, you agree to first attempt resolution by contacting us at **damr58h@gmail.com**. If the dispute cannot be resolved within 30 days, it shall be submitted to the competent courts of the Czech Republic. For consumer disputes, the court of the consumer's domicile shall have jurisdiction.

17. Severability

If any provision of these Terms is found to be unenforceable or invalid, that provision shall be limited or eliminated to the minimum extent necessary, and the remaining provisions shall remain in full force and effect.

18. Non-Reliance on AI Output

You acknowledge and agree that:

- (a) AI-extracted data and chat responses are provided for informational purposes only;
- (b) OmniParse does not warrant accuracy, completeness, timeliness, or fitness for any specific purpose;
- (c) You are solely responsible for verifying all extracted data and chat responses against source documents before any use;
- (d) If you rely on inaccurate AI output for financial decisions, tax filings, legal actions, vendor payments, or any other purpose, OmniParse is not liable for resulting damages, losses, or penalties;
- (e) The confidence scores provided are model-generated estimates only and do not guarantee accuracy;
- (f) Chat responses may contain errors, hallucinations, or be out of context — always cross-reference against the underlying invoice records before acting on them;
- (g) OmniParse may modify, suspend, or discontinue AI features or change AI providers, with reasonable notice for material changes as described in § 7b.

By accepting these Terms, you agree that you will not hold OmniParse, its operator, or its AI providers liable for losses arising from reliance on AI output.

20. Force Majeure

Neither party shall be liable for any failure or delay in performing its obligations under these Terms (except for payment obligations) to the extent such failure or delay is caused by events beyond its reasonable control, including but not limited to acts of God, natural disasters, war, terrorism, civil unrest, epidemics or pandemics, governmental actions, labor disputes, internet or telecommunications failures, or failures of third-party AI providers, hosting providers, or payment processors. The affected party shall notify the other party without undue delay and shall use commercially reasonable efforts to resume performance as soon as reasonably practicable.

21. Notices

Notices from OmniParse to you: Will be sent to your account email address or displayed in the Service. Email notices are deemed received 24 hours after sending. In-Service notices are deemed received upon first display.

Notices from you to OmniParse: Send to damr58h@gmail.com with a clear subject line (e.g., "DSR", "DPA", "Refund Request", "Account Termination"). Legal notices are deemed received on the date we reply acknowledging receipt.

22. Acceptable Use Policy

You agree not to:

- Upload documents you do not have the right to process (copyright, privacy, or other legal restrictions)
- Upload documents containing personal data of EU data subjects without a valid legal basis for transfer to US-based AI providers;
- Use automated scraping, bots, or rate-limit bypass tools to overwhelm the Service;
- Attempt IDOR (insecure direct object reference), prompt injection, SQL injection, XSS, or other exploits against the Service;
- Upload child sexual abuse material (CSAM), hate speech, content promoting violence or terrorism, or content that is illegal in your jurisdiction or the operator's jurisdiction (Czech Republic)
- Use the Service for fraud, money laundering, tax evasion, or any other illegal activity;
- Reverse-engineer, decompile, or attempt to extract AI model weights, training data, or proprietary algorithms;
- Resell or sublicense access to the Service without written permission;
- Interfere with the proper functioning of the Service, including by introducing viruses, malware, or other malicious code.
- Create multiple accounts to bypass plan limits or rate limits;
- Use the Service to burn AI provider quotas in an abusive or automated manner.

Account monitoring and enforcement: We monitor account activity to detect potential abuse, including excessive API usage, rapid account creation, and unusual patterns of document processing. We reserve the right to:

- **Freeze** accounts suspected of abuse — the user can still log in but AI features are blocked;
- **Suspend** accounts suspected of abuse, with prior written notice by email where feasible. For paid accounts, we will provide at least 7 days' notice before suspension unless the violation is urgent (e.g., CSAM, ongoing

fraud

- **Delete** accounts that materially violate this Acceptable Use Policy, including all associated data. For paid accounts, you may appeal a deletion decision within 14 days by emailing damr58h@gmail.com;
- **Report** suspected illegal activity to the relevant authorities (Czech Police, ÚOOÚ, Europol).

If your account is frozen, you will see a message indicating the reason when you attempt to use AI features. To request reinstatement, contact damr58h@gmail.com.

23. Contact

For questions about these Terms: damr58h@gmail.com

Compliant with EU Directive 93/13/EEC (unfair contract terms), EU Directive 2011/83/EU (consumer rights, including right of withdrawal), EU Directive (EU) 2023/2673 (one-click withdrawal button, effective 19 June 2026), Czech Act No. 89/2012 Coll. (Civil Code, including §§ 419, 1752, 1753, 1829, 1837, 2913), Czech Act No. 634/1992 Coll. (Consumer Protection, including §§ 16, 20b–20e), Czech Act No. 110/2019 Coll. (Czech GDPR implementation), Czech Act No. 127/2005 Coll. § 89 (Electronic Communications — ePrivacy implementation), and Czech Act No. 480/2004 Coll. (Information Society Services — unsolicited communications).

Privacy Policy

Last updated: September 18, 2026

1. Data Controller

The data controller for personal data processed through this Service is **Simon Curda** (natural person), based in the Czech Republic. For GDPR purposes (EU Regulation 2016/679), the Controller is responsible for ensuring compliance with data protection law. The Controller has not appointed a Data Protection Officer (DPO) as none is legally required under GDPR Art. 37 for an operation of this scope and nature; all GDPR-related matters are handled directly by the Controller.

Contact: damr58h@gmail.com (use this for all GDPR-related requests: access, rectification, erasure, restriction, portability, objection, and consent withdrawal). We aim to respond within 30 days as required by GDPR Art. 12(3). Where a request is complex or numerous, this period may be extended by a further two months pursuant to GDPR Art. 12(3), in which case we will inform you of the extension and the reasons within the initial 30-day period.

Identity of the Controller and contact details (GDPR Art. 13(1)(a)–(b)): Simon Curda, natural person, domicile Czech Republic. Email: damr58h@gmail.com. For Data Subject Requests, please use the subject line "DSR — [your name]". For B2B / DPA matters, please use the subject line "DPA — [company name]".

2. Personal Data We Collect

Account information: When you register, we collect your name and email address. Your password is hashed using bcrypt (12 rounds) and never stored in plaintext.

Documents you upload: When you use the document parsing feature, your files (PDF, JPEG, PNG, WebP) are sent to AI processing providers (see Section 6 for the full list). The original file is stored in your account database (base64-encoded) for 30 days so you can preview it in the dashboard, then deleted from the active environment after 30 days. The extracted data is stored in your account until you delete it or terminate your account.

Chat history: Messages you send to the AI chat assistant and the responses returned are stored in your account database until you delete them via the dashboard's Clear button or by deleting your account.

Technical data: We collect standard server logs including your IP address, browser type, and request timestamps for security and service stability purposes.

Local storage data: The application uses your browser's local storage to remember your authentication token, theme preference, and cookie consent choice. Invoice and chat data is stored server-side in your account, not in local storage.

3. Legal Basis for Processing (GDPR Art. 6)

- **Account creation and authentication:** Contractual necessity (Art. 6(1)(b)) — necessary to provide the Service you requested.
- **Document parsing:** Consent (Art. 6(1)(a)) — you actively upload documents for processing. You can withdraw this consent by deleting your invoices.
- **Chat interactions:** Consent (Art. 6(1)(a)) — you actively send messages to the AI chat.
- **Terms acceptance:** Consent (Art. 6(1)(a)) — recorded at signup with timestamps.
- **Security, service stability, and abuse prevention:** Legitimate interest (Art. 6(1)(f)) — server logs for abuse prevention, rate limiting, fraud detection, and account monitoring. We monitor account activity patterns (invoice volume, chat usage, email scanning frequency, account age) to detect potential abuse. Accounts suspected of abuse may be frozen (AI features blocked) or deleted in accordance with our Terms of Service. A legitimate-interest balancing test has been conducted and is documented in our internal Records of Processing Activities (ROPA it is available to supervisory authorities on request).
- **Local storage preferences:** Consent via cookie banner (Art. 6(1)(a) + ePrivacy Directive Art. 5(3), implemented in Czech Act No. 127/2005 Coll. § 89).
- **Processing of personal data appearing in customer-uploaded invoices** (e.g., names of vendor employees, contact persons, signatories): Legitimate interest (Art. 6(1)(f)) — necessary to provide the document-extraction Service the customer requested. Where such data reveals special categories (Art. 9), processing is limited to what is strictly necessary for invoice extraction and is not used for any other purpose.

4. Data Retention (GDPR Art. 5(1)(e))

We retain personal data only for as long as necessary for the purposes described in this Policy, or as required by applicable law. The following retention periods apply:

Data category Retention period Basis / exception **Uploaded document files** (PDFs, images, base64-encoded in primary DB) 30 days in active environment Auto-purged after 30 days. Extraction results remain in account.

Extracted invoice data Until account deletion (or user deletes individual invoice) Deleted within 90 days of account termination. **Chat history** Until cleared by user or account deletion Cascade-deleted with account. **Audit logs** (invoice edits, approvals, deletions) Lifetime of account Cascade-deleted with account. **Account data** (email,

*name, hashed password, settings) Until account deletion User can delete at any time from Settings. **Server logs** (Vercel — IP, browser, timestamps) 30 days Vercel Hobby/Pro plans auto-delete after 30 days. **Database backups** (Supabase PITR) 7 days (rolling) Disaster recovery only; overwritten on rolling basis. **Stripe billing records** Per Stripe policy (~10 years) Required by financial/tax regulations. Governed by Stripe's privacy policy. **Incident logs** (security incidents involving AI systems) 2 years Best practice for regulatory audit purposes (AI Act Art. 12). **Withdrawal / consent records** (termsAcceptedAt, withdrawalAcknowledgedAt, withdrawnAt, lastTosEmailSentAt) Lifetime of account + 90 days Evidence of consent for regulatory audit; cascade-deleted with account.*

No indefinite retention. We do not retain personal data indefinitely. Where longer retention is required by law (e.g., financial records under Czech accounting regulations), retention is limited to the statutory period and data is then deleted.

Relationship between 30-day and 90-day periods: Uploaded document *files* (the binary PDF/image) are purged after 30 days. *Extracted data* (the structured JSON — vendor, invoice number, line items) remains in your account for as long as your account is active, and is deleted within 90 days of account termination. These are different data categories with different retention needs; the 30-day period applies to file binaries, the 90-day period applies to post-termination cleanup of all remaining personal data.

5. Your Rights Under GDPR (Art. 12-23)

You have the following rights regarding your personal data:

- **Right of access (Art. 15):** Request a copy of all personal data we hold about you. Use the "Export my data" feature in Settings, or contact us.
- **Right to rectification (Art. 16):** Request correction of inaccurate personal data. Edit your profile from Settings.
- **Right to erasure / "right to be forgotten" (Art. 17):** Request deletion of your personal data. Use the "Delete account" feature in Settings — this deletes all your data within seconds.
- **Right to restriction of processing (Art. 18):** Request that we limit how we use your data. Contact us to request this.
- **Right to data portability (Art. 20):** Receive your data in a structured, machine-readable format. Use the "Export my data" feature in Settings.
- **Right to object (Art. 21):** Object to processing based on legitimate interests (e.g. server log collection).
- **Right to withdraw consent (Art. 7(3)):** Withdraw consent at any time via cookie settings, by deleting your invoices and chat history, or by contacting us.
- **Rights related to automated decision-making and profiling (Art. 22):** OmniParse does **not** engage in solely automated decision-making producing legal or similarly significant effects concerning you. AI extraction and chat responses are tools to assist you; all material decisions (e.g., invoice approval, payment authorization) require your explicit human action. You have the right not to be subject to a decision based solely on automated processing; if you believe any such decision has occurred, contact us immediately.

To exercise any of these rights, contact us at damr58h@gmail.com. We will respond within 30 days as required by GDPR Art. 12(3). You also have the right to lodge a complaint with a supervisory authority — in particular the Czech Office for Personal Data Protection (ÚOOÚ) at [uoou.cz](https://www.uoou.cz) (<https://www.uoou.cz>), or, if you reside in another EU/EEA Member State, the supervisory authority of your habitual residence.

6. Subprocessors and AI Providers

We use the following third-party services to provide the Service. Each processes personal data on our behalf as a subprocessor under GDPR Art. 28. Data Processing Agreements (DPAs) and Standard Contractual Clauses (SCCs) are in place with all *active* providers (see per-provider status below).

EU-based AI provider available. Mistral AI (Paris, France) processes documents within the EU. Transfers to Mistral are expected to remain within the EEA, subject to Mistral's applicable terms. We prioritize Mistral for vision extraction and chat when available. When `MISTRAL_DISABLE_TRAINING=true` is set in our server configuration (default), we send `usage_options={"enable_training": false}` on every Mistral API call to disable training on prompt content (effective on paid Mistral tier; on free tier, Mistral may ignore or reject this parameter).

- **Vercel Inc.** (United States, DPF-certified) — Web hosting and serverless function execution. Processes: IP addresses, request metadata. **Vercel DPA** (<https://vercel.com/legal/dpa>).
- **Supabase Inc.** (Ireland, EU) — Primary database hosting (PostgreSQL). Processes: all stored account data, invoices, chat messages. DPA built into Terms of Service (effective August 1, 2026) — automatically applies to all customers, no separate signature required. **Supabase DPA** (<https://supabase.com/legal/customer-resources/data-processing-addendum>). Subprocessor change notifications: supabase.com/legal/customer-resources/subprocessor-list (<https://supabase.com/legal/customer-resources/subprocessor-list>).
- **Mistral AI** (Paris, France — EU) — Primary AI inference provider. Processes: uploaded document content (transiently) for vision extraction, chat messages (transiently). **EU-based — no SCC required.** Training opt-out (`usage_options.enable_training=false`) is sent on every request when `MISTRAL_DISABLE_TRAINING=true` (default). **Mistral Privacy Policy** (<https://mistral.ai/legal/privacy-policy>).
- **Groq Inc.** (United States) — Secondary AI inference provider (after Mistral). Processes: uploaded document content (transiently), chat messages (transiently). **SCCs confirmed** (EU SCC Module 2, self-executing upon acceptance of Groq Services Agreement; DPA dated October 15, 2025). DPA accessible at **console.groq.com** (<https://console.groq.com>). Governing law: Ireland. *Note: Groq is optional in our configuration — operators may deploy in EU-only mode (Mistral only) by leaving `GROQ_API_KEY` unset.*
- **Google LLC** (United States, DPF-certified) — Final fallback AI inference provider. Models used: Gemini 2.0 Flash, Gemini 2.5 Flash, Gemini 1.5 Flash. **DISABLED BY DEFAULT** — uses Google's AI Studio free-tier endpoint (`generative_language.googleapis.com`), whose data-handling terms are weaker than Google Cloud Vertex AI (which has a self-executing DPA). Operator must set `ENABLE_GOOGLE_GEMINI=true` after reviewing Google AI Studio terms. For production with EU personal data, operators are strongly encouraged to migrate to Vertex AI instead. Google LLC itself is DPF-certified. **Google Cloud DPA** (<https://cloud.google.com/terms/data-processing-addendum>) (applies to Vertex AI path, not AI Studio free tier).
- **Stripe Inc.** (United States, DPF-certified) — Payment processing for paid plans. Processes: email, billing details. **Stripe DPA** (<https://stripe.com/legal/dpa>).

OpenRouter — DISABLED by default. OpenRouter is an AI inference aggregator that routes requests to underlying model providers (Meta, Nvidia, Google, etc.). OpenRouter's free-tier models (all model IDs ending in `:free`) typically permit training on prompt content, and there is no DPA or SCC in place with OpenRouter as of September 2026. **OpenRouter is therefore disabled by default in our production configuration** (the `ENABLE_OPENROUTER` environment variable must be explicitly set to `true` by the operator after completing

their own DPA/SCC review). When disabled, OpenRouter is never called — document content and chat messages are never transmitted to OpenRouter or its underlying providers.

AI vision cascade (production default): Mistral (EU, primary) → Groq (US, SCCs confirmed, secondary). OpenRouter and Google Gemini are **disabled by default** and only used if the operator explicitly enables them after completing their own DPA/SCC review (OpenRouter) or AI Studio terms review (Google Gemini).

SCC Status (as of September 18, 2026): **Supabase Inc.** (Ireland, EU) — DPA built into Terms of Service (effective August 1, 2026). Automatically applies to all customers. EU-based — no SCC required. **Vercel Inc.** (United States) — DPF-certified. DPA available at vercel.com/legal/dpa. **Groq Inc.** — SCCs confirmed in effect (DPA dated October 15, 2025; EU SCC Module 2 self-executing upon acceptance of Groq Services Agreement). DPA accessible at **console.groq.com** (<https://console.groq.com>). Governing law: Ireland. Competent authority: Irish DPC. 72-hour breach notification. DPF certification status pending verification. **Mistral AI** — EU-based (Paris, France). No Chapter V transfer issue; no SCC required. Training opt-out sent on every request (MISTRAL_DISABLE_TRAINING=true). **Google LLC (Gemini)** — **DISABLED BY DEFAULT**. Uses AI Studio free-tier endpoint (generativelanguage.googleapis.com). Google Cloud DPA (cloud.google.com/terms/data-processing-addendum) applies to the Vertex AI path, NOT the AI Studio free tier. Google LLC is DPF-certified. Operator must set ENABLE_GOOGLE_GEMINI=true after reviewing AI Studio terms. For EU personal data, migrate to Vertex AI. **Stripe Inc.** (United States) — DPF-certified. DPA available at stripe.com/legal/dpa. **OpenRouter** — **DISABLED by default**. Not used in production unless operator explicitly sets ENABLE_OPENROUTER=true after completing DPA/SCC review. When disabled, no data is transferred to OpenRouter or its underlying providers.

EU users: AI processing via Mistral (EU-based) and Groq (US-based, SCCs confirmed) has appropriate safeguards in place under GDPR. Google is DPF-certified. OpenRouter is disabled by default. However, users should note that documents uploaded to the Service may contain personal data of multiple data subjects (e.g. vendor names, employee names, email addresses, bank details). The user uploading such documents is responsible for ensuring they have a valid legal basis under GDPR Art. 6 for processing that data, and that the transfer to AI providers is lawful under Chapter V.

Model training — production configuration: OmniParse sends `usage_options={"enable_training": false}` on every Mistral API call (default server configuration) to disable training on prompt content. This opt-out is effective on paid Mistral tier; on free Mistral tier, Mistral may ignore or reject this parameter. OpenRouter (which has no training opt-out on free-tier models) is disabled by default. Google Gemini (AI Studio free tier, whose data-handling terms are weaker than Google Cloud Vertex AI) is also disabled by default. Groq's data-handling terms are governed by the Groq DPA and SCCs. **OmniParse itself does not route customer content through AI API tiers that permit provider training on customer content** when the default configuration (Mistral with training opt-out, OpenRouter disabled, Google Gemini disabled) is in effect. Operators who change this configuration must update this Policy accordingly.

Important — Schrems III risk: The EU-US Data Privacy Framework (DPF) adequacy decision (10 July 2023) is currently subject to a CJEU appeal ("Schrems III"). If the DPF is invalidated, transfers to DPF-certified US recipients (Vercel, Google, Stripe) would need to rely on SCCs + Transfer Impact Assessment (TIA) instead. We maintain SCCs with Groq as a contingency and will execute SCCs with other US providers if the DPF is invalidated.

7. Security Measures (GDPR Art. 32)

OmniParse implements appropriate technical and organizational measures to ensure a level of security appropriate to the risk, taking into account the state of the art and the costs of implementation. Measures include:

- **Encryption in transit:** HTTPS/TLS 1.2+ for all connections (Vercel-managed certificates, auto-renewed).
- **Encryption at rest:** Supabase (PostgreSQL) provides transparent disk encryption (TDE) on managed infrastructure. IMAP credentials stored by users are encrypted with AES-256-GCM.
- **Password storage:** Passwords are hashed with bcrypt (12 rounds) and never stored in plaintext or reversible form.
- **Authentication:** JWT-based authentication with 7-day token expiry. Tokens are signed with a server-side secret and rotated periodically.
- **Rate limiting:** 5 login attempts per minute per IP on authentication endpoints; broader rate limits on AI endpoints to prevent abuse.
- **Security headers:** Content-Security-Policy, Strict-Transport-Security (HSTS, 2-year max-age, preload), X-Frame-Options: DENY, X-Content-Type-Options: nosniff, Referrer-Policy: strict-origin-when-cross-origin.
- **Per-user data isolation:** Every database query includes a userId filter — no insecure direct object reference (IDOR) is possible at the query layer.
- **Prompt-injection defense:** All chat messages are screened against 30+ prompt-injection attack patterns before being sent to AI providers.
- **Account deletion:** Transactional cascade delete ensures full erasure of all user-owned rows, including invoices, chat, audit logs, and IMAP credentials.
- **Vendor due diligence:** Subprocessors are assessed before engagement and re-assessed at least annually. DPAs and (where applicable) SCCs are in place (see §6 and §11).
- **Incident response:** A documented incident-response procedure exists. Personal-data breaches are notified to the Czech DPA (ÚOOÚ) within 72 hours of becoming aware, where required by GDPR Art. 33, and to affected data subjects without undue delay where required by GDPR Art. 34.

8. Children's Data (GDPR Art. 8)

The Service is an invoice-extraction and analytics platform intended for use by adults acting in a professional or business capacity. The Service is **not directed at children**, and we do not knowingly collect personal data from children. In the Czech Republic, the age of digital consent under GDPR Art. 8 (as implemented by Act No. 110/2019 Coll.) is **15**. In other EU/EEA Member States the age ranges from 13 to 16 depending on national law.

Users must be at least **15 years old** to register an account. By registering, you confirm that you have reached at least the age of digital consent in your jurisdiction (and where the local age is higher than 15, that you meet that higher age requirement). Users aged between the applicable digital-consent age and 18 must additionally obtain parental or guardian consent to use the Service and to enter into these Terms.

If we learn that we have collected personal data from a child under the applicable age of digital consent without verifiable parental consent, we will delete that data promptly. Contact us at damr58h@gmail.com if you believe a child has provided us with personal data.

9. Czech Republic Compliance

For users in the Czech Republic, the following applies in addition to GDPR:

- **Act No. 110/2019 Coll.** on the processing of personal data — Czech implementation of GDPR. Sets the age of digital consent at 15 (Art. 8 GDPR) and codifies national derogations.
- **Act No. 89/2012 Coll.** (Civil Code) — contractual relationships, consumer protection, liability regime (incl. § 2913 on limitation of liability).
- **Act No. 127/2005 Coll.** (Electronic Communications Act, "Lex Telecom") § 89 — implements the ePrivacy Directive (2002/58/EC) cookie/storage consent rule, as amended effective 1 January 2022. Requires opt-in consent before storing or accessing information on a user's terminal equipment, except where strictly necessary (Art. 5(3) exception).
- **Act No. 480/2004 Coll.** (Information Society Services Act) — governs unsolicited commercial communications (spam) and certain information-society-service provider obligations. *Note: this Act is not the cookie-consent statute; that role belongs to Act No. 127/2005 Coll. § 89.*
- **Act No. 634/1992 Coll.** (Consumer Protection Act) §§ 16, 20b–20e — out-of-court dispute resolution (ADR) before Česká obchodní inspekce (ČOI).
- **Digital Services Act (Regulation (EU) 2022/2065):** Directly applicable in the Czech Republic since 17 February 2024. OmniParse operates as a hosting service (within the meaning of Art. 6 DSA) for customer-uploaded documents; it does not qualify as an online platform and is therefore not subject to the additional obligations in Chapter IV DSA.
- **Supervisory authority (data protection):** Úřad pro ochranu osobních údajů (Office for Personal Data Protection), Pplk. Sochora 27, 170 00 Praha 7, Czech Republic. Website: [uoou.cz](https://www.uoou.cz) (<https://www.uoou.cz>).
- **Consumer-protection authority:** Česká obchodní inspekce (ČOI), Štěpánská 567, 120 00 Praha 2. Website: coi.gov.cz (<https://coi.gov.cz>).

10. Canadian PIPEDA Compliance

For users accessing the Service from Canada, OmniParse complies with the Personal Information Protection and Electronic Documents Act (PIPEDA, R.S.C. 2000, c. P-8.9). Bill C-27 (which would have replaced PIPEDA with the Consumer Privacy Protection Act) died on the Order Paper on 6 January 2025 and did not become law; PIPEDA therefore remains in force unamended.

The ten fair-information principles in Schedule 1 of PIPEDA are respected, in particular:

- **Accountability:** Simon Curda (natural person, Czech Republic) is responsible for personal information under PIPEDA and has designated damr58h@gmail.com as the contact for privacy matters.
- **Identifying purposes:** Purposes are described in §2 of this Policy.
- **Consent:** Consent is obtained at signup (Terms acceptance) and via the cookie banner for non-essential storage. Canadian users may withdraw consent at any time by contacting us or by deleting their account.
- **Limiting collection:** We collect only what is necessary for the Service (see §2).
- **Limiting use, disclosure, retention:** Retention periods in §4 apply.
- **Accuracy:** Users can correct their data via Settings; AI-extracted data can be edited in the dashboard.
- **Safeguards:** See §7 (security measures).
- **Openness:** This Policy and the Cookie Policy are publicly available without account.

- **Individual access:** Canadian users may request access to, correction of, or deletion of their personal data by emailing damr58h@gmail.com.
- **Challenging compliance:** Canadian users may challenge our compliance by contacting us or by filing a complaint with the Office of the Privacy Commissioner of Canada at [priv.gc.ca](https://www.priv.gc.ca) (<https://www.priv.gc.ca>).

Mandatory breach notification (PIPEDA § 10.1 and Breach of Security Safeguards Regulations): Where a breach of security safeguards involving personal information under our control creates a real risk of significant harm to a Canadian individual, we will notify the affected individual and the Privacy Commissioner of Canada as soon as feasible after determining the breach has occurred, and we will keep a record of the breach for at least 24 months.

11. Data Processing Addendum (DPA)

This Data Processing Addendum ("DPA") applies when you (the "Customer") act as a data controller and use the Service to process personal data on behalf of data subjects. This DPA reflects the parties' agreement with respect to the processing of personal data under GDPR Art. 28.

11.1 Roles and Scope

- **Customer** is the data controller. Customer determines the purposes and means of processing.
- **OmniParse** (operated by Simon Curda) is the data processor. We process personal data only on documented instructions from Customer.
- We process documents and extracted data per your instructions (via the Service's UI and API). We do not use your data for any purpose other than providing the Service.

11.2 Subprocessors

OmniParse engages the following subprocessors. Customer grants general written authorization for OmniParse to engage these subprocessors; the current list is maintained in this Privacy Policy §6 and material changes will be notified 30 days in advance per §11.7:

- **Vercel Inc.** (United States, DPF-certified) — Web hosting and serverless function execution
- **Supabase Inc.** (Ireland, EU) — Primary database hosting (PostgreSQL)
- **Mistral AI** (Paris, France — EU) — Primary AI inference (vision + text). Training opt-out sent on every request.
- **Groq Inc.** (United States, SCCs confirmed) — AI inference (secondary fallback). Optional — operator may disable for EU-only mode.
- **Google LLC** (United States, DPF-certified) — AI inference (final fallback, Gemini models)
- **Stripe Inc.** (United States, DPF-certified) — Payment processing (paid plans only)
- **OpenRouter** (United States) — **DISABLED BY DEFAULT.** Not engaged unless operator explicitly sets `ENABLE_OPENROUTER=true` after completing DPA/SCC review. When disabled, no data is transferred to OpenRouter.

11.3 Data Subject Rights

Customer is responsible for responding to data subject rights requests (access, rectification, erasure, portability, objection). OmniParse will assist Customer with such requests, including by exporting or deleting personal data upon Customer's written request. Contact damr58h@gmail.com with the subject "DPA — Data Subject Request".

11.4 Security Measures (GDPR Art. 32)

OmniParse implements the following technical and organizational measures (cross-referenced from Privacy Policy §7):

- Encryption in transit (HTTPS/TLS 1.2+) for all connections
- Encryption at rest: Supabase-managed transparent disk encryption (TDE AES-256-GCM for IMAP credentials at rest)
- bcrypt password hashing (12 rounds)
- JWT-based authentication with 7-day expiry
- Per-request user scoping — no query returns cross-user data
- File data auto-purge after 30 days
- Account deletion with cascade (full erasure)
- Audit log of significant actions (invoice create/approve/delete)
- Rate limiting on auth endpoints (5 attempts/min/IP)
- Security headers: CSP, HSTS (2-year preload), X-Frame-Options: DENY, X-Content-Type-Options: nosniff, Referrer-Policy
- Prompt-injection screening on all chat messages (30+ patterns)
- Vendor due diligence: subprocessors assessed before engagement and re-assessed at least annually

11.5 International Transfers (GDPR Chapter V)

Personal data may be transferred to the United States for AI processing by Groq and Google. The EU-US Data Privacy Framework (DPF) adequacy decision of 10 July 2023 remains in force as of the date of this Policy, although it is subject to an ongoing CJEU appeal ("Schrems III"). For US recipients self-certified under the DPF, the DPF alone provides a valid transfer mechanism and SCCs are not additionally required. For US recipients NOT DPF-certified, Standard Contractual Clauses (SCCs) consistent with the Schrems II ruling, together with a Transfer Impact Assessment (TIA), are required. **OpenRouter is DISABLED by default** — no personal data is transferred to OpenRouter unless the operator explicitly enables it after completing DPA/SCC review. **SCC status by provider (as of September 18, 2026):**

- **Groq Inc.** — SCCs confirmed in effect (DPA dated October 15, 2025; EU SCC Module 2 self-executing upon acceptance of Groq Services Agreement). DPA accessible at console.groq.com (<https://console.groq.com>). Governing law: Ireland. Competent authority: Irish Data Protection Commission. 72-hour breach notification. Groq's DPF certification status is pending verification at [dataprivacyframework.gov](https://www.dataprivacyframework.gov) (<https://www.dataprivacyframework.gov>); SCCs remain in place as a safeguard.
- **Mistral AI** — EU-based (Paris, France). No Chapter V transfer issue; no SCC required. Training opt-out (`usage_options.enable_training=false`) sent on every request when `MISTRAL_DISABLE_TRAINING=true` (default).
- **Google LLC** — Google Cloud DPA self-executing upon acceptance of Google Cloud Terms; available at cloud.google.com/terms/data-processing-addendum (<https://cloud.google.com/terms/data-processing-addendum>). Google LLC is DPF-certified (verified September 2026 SCCs serve as fallback should the DPF be invalidated by Schrems III).
- **Vercel Inc.** — Vercel is DPF-certified. DPA available at vercel.com/legal/dpa (<https://vercel.com/legal/dpa>).

- **Stripe Inc.** — Stripe is DPF-certified. DPA available at stripe.com/legal/dpa (<https://stripe.com/legal/dpa>).
- **OpenRouter** — **DISABLED BY DEFAULT.** No personal data is transferred to OpenRouter unless the operator explicitly sets `ENABLE_OPENROUTER=true`. Before enabling, the operator must complete a DPA with OpenRouter, verify SCCs or DPF certification, switch to paid-tier models (to disable training), and update this Policy.

OmniParse has confirmed SCCs with Groq (US-based) and uses Mistral (EU-based) as the primary provider. The cascade prioritizes Mistral first, then Groq — both have appropriate safeguards. Google is covered by the DPF. OpenRouter is disabled by default and not used unless the operator explicitly enables it. Users should be aware that documents may contain personal data of multiple data subjects — the user is responsible for ensuring they have a valid legal basis for processing and transferring such data under GDPR Art. 6 and Chapter V.

11.6 Duration and Deletion

This DPA continues for the duration of the Customer's use of the Service. Upon account termination, file binaries are deleted within 30 days, and all other personal data (extracted data, audit logs, chat history, IMAP credentials) is deleted within 90 days, except where longer retention is required by law (e.g., financial records under Czech accounting regulations). Backup copies maintained by Supabase for disaster recovery are deleted within 7 days on a rolling basis. This is consistent with the retention schedule in §4 of this Privacy Policy.

11.7 Subprocessor Authorization

Customer grants general written authorization for OmniParse to engage the subprocessors listed in Section 11.2. OmniParse will notify Customer of any planned changes to the subprocessor list at least 30 days in advance via the Service or by email. Customer may object to a new subprocessor in writing within 30 days of notification. In the event of a reasonable objection, OmniParse will use commercially reasonable efforts to make available an alternative or recommend a change to Customer's use of the Service to avoid processing by the objected-to subprocessor. If OmniParse is unable to provide a reasonable alternative, Customer may suspend or terminate the affected Service.

11.9 Personal Data Breach Notification (GDPR Art. 33-34)

OmniParse will notify Customer without undue delay after becoming aware of a personal-data breach affecting Customer's personal data. Notification will include (i) a description of the nature of the breach, (ii) the likely consequences, (iii) the measures taken or proposed, and (iv) the contact point for further information. Where the breach is likely to result in a risk to the rights and freedoms of natural persons, OmniParse will assist Customer in fulfilling Customer's obligation under GDPR Art. 33 to notify the supervisory authority within 72 hours, and under Art. 34 to notify affected data subjects without undue delay.

11.10 Audit Rights

Customer may audit OmniParse's compliance with this DPA, upon reasonable prior written notice (at least 30 days) and not more than once per calendar year, during normal business hours and in a manner that does not interfere with OmniParse's operations. Alternatively, OmniParse may provide a current third-party audit report (e.g., SOC 2 Type II, ISO 27001) or equivalent documentation in lieu of an on-site audit, which the parties agree is a reasonable alternative under GDPR Art. 28(3)(h).

11.11 Acceptance and Activation

For B2B customers (businesses, organizations, accounting firms): This DPA takes effect and is binding between OmniParse (as processor) and the B2B Customer (as controller) upon the later of: (a) the Customer's acceptance of the Terms of Service, and (b) the Customer's first use of the Service to process personal data. The DPA is incorporated by reference into the Terms of Service; no separate signature is required for it to take effect (per GDPR Art. 28(9) and Recital 81, which permit DPAs to be formed by written agreement including electronic form).

Activation mechanism: B2B Customers who require a standalone signed DPA (e.g., for their own compliance records, customer-facing DPAs, or large Enterprise contracts) may request one by contacting damr58h@gmail.com with the subject line "DPA Activation — [Company Name]". We will execute a standalone Data Processing Agreement within 10 business days. Until then, this embedded DPA (Privacy Policy §11) governs the processing relationship.

For individual users (consumers): By uploading documents to the Service and checking the data-transfer consent box in the Upload tab, you acknowledge that this DPA applies to the extent you act as a data controller processing personal data through the Service. For consumers acting in a personal capacity (not processing data on behalf of an organization), the controller-processor relationship under GDPR Art. 28 does not strictly arise — you are the controller and OmniParse is your processor under the Terms of Service.

Version control: The "Last updated" date at the top of this Privacy Policy reflects the current DPA version. Material changes to this DPA will be notified to B2B Customers at least 30 days in advance per §11.7 (Subprocessor Authorization mechanism). Prior versions are retained and available on request.

12. Privacy Notice for Invoice Data Subjects (GDPR Art. 14)

Where customers upload documents (e.g., invoices) that contain personal data of individuals other than themselves — for example, contact persons at vendor companies, employee names, signatories, or bank-account holders — that personal data is processed by OmniParse on the customer's instructions. The customer is the controller with respect to that data; OmniParse acts as a processor on the customer's behalf.

Pursuant to GDPR Art. 14, and in reliance on the Art. 14(5)(b) exemption (providing the information would involve a disproportionate effort because OmniParse has no direct relationship with those data subjects and they may be numerous), this public Privacy Notice serves as the transparency mechanism. We ask our customers to inform their data subjects of the use of OmniParse where feasible and to direct them to this Notice.

If your personal data appears in a document processed by OmniParse:

- **Who is processing your data:** The OmniParse customer who uploaded the document is the data controller. OmniParse (Simon Curda, natural person, Czech Republic) acts as a processor on their behalf.
- **Purposes:** Extraction of structured data from invoices for accounting, accounts-payable, and vendor-management purposes.
- **Legal basis:** Legitimate interest of the customer in efficient invoice processing (GDPR Art. 6(1)(f)).
- **Recipients:** AI inference providers (Mistral, Groq, OpenRouter, Google) and infrastructure providers (Vercel, Supabase, Stripe) — see §6 above.
- **Retention:** Personal data is deleted within 90 days of the customer's account termination, or earlier if the customer deletes the document.
- **Your rights:** You have rights of access, rectification, erasure, restriction, portability, and objection. To exercise these rights, contact the customer (the controller) directly; if you cannot reach them, contact us at

damr58h@gmail.com and we will route your request.

- **Right to complain:** You may lodge a complaint with the Czech DPA (ÚOOÚ) or with the supervisory authority of your habitual residence.

13. Changes to This Policy

We may update this Privacy Policy from time to time. Material changes will be notified via the Service or by email at least 30 days before taking effect. The "Last updated" date at the top of this page indicates when this Policy was last revised. Continued use of the Service after changes constitutes acceptance of the updated Policy.

14. Contact

- **Data controller:** Simon Curda, natural person, Czech Republic
- **General contact & all GDPR requests:** damr58h@gmail.com
- **B2B / DPA matters:** damr58h@gmail.com (subject line "DPA — [company name]")
- **Data Subject Requests (Art. 15-22):** damr58h@gmail.com (subject line "DSR — [your name]")
- **Czech supervisory authority:** uoou.cz (<https://www.uoou.cz>)
- **Canadian Commissioner:** priv.gc.ca (<https://www.priv.gc.ca>)

Cookie Policy

Last updated: September 18, 2026

1. What Are Cookies and Local Storage

Cookies are small text files stored on your device by websites you visit. OmniParse primarily uses browser Local Storage rather than traditional cookies. Under the EU ePrivacy Directive (2002/58/EC) as amended by Directive 2009/136/EC, the consent requirements apply to both cookies and Local Storage. Each item stored is classified as either essential or non-essential based on its purpose.

2. What We Store and Why

Essential storage (always active — exempt from consent per ePrivacy Art. 5(3), as implemented by Czech Act No. 127/2005 Coll. § 89):

- **Authentication token (op_token):** Stores your JWT session token in Local Storage so you stay signed in across page refreshes. *Retention:* 7 days (matching JWT expiry), then auto-cleared on next page load. *Purpose:* Strictly necessary for the Service to function.
- **Cookie consent record (omniparse_cookie_consent):** Stores your consent choice and timestamp. *Retention:* 12 months from last update, then re-prompted. *Purpose:* Prevents re-showing the consent banner and records your consent decision for accountability.

Non-essential storage (requires consent — only set after you accept via the cookie banner):

- **Theme preference (theme):** Stores your light/dark mode choice. *Retention:* Until withdrawn or browser data cleared. *Purpose:* User preference (UI customization).

- **Keyboard shortcuts (op_shortcuts):** Stores custom keyboard shortcut assignments. *Retention:* Until withdrawn or browser data cleared. *Purpose:* User preference (UI customization).
- **Legal consent (op_legal_consent):** Records GDPR data-transfer consent for document uploads. *Retention:* Duration of account; refreshed on each upload if revoked. *Purpose:* Compliance (GDPR consent evidence).
- **Crash logs (op_crash_log):** Stores recent client-side errors for debugging. *Retention:* Rolling window — last 10 errors only, FIFO overwrite. *Purpose:* Service stability and debugging. No personal data is logged beyond error messages.

Analytics:

- Currently, no analytics cookies or storage items are active. If we implement analytics in the future, we will (a) update this Policy with the specific item name, purpose, and retention; (b) seek your explicit consent via the cookie banner; and (c) ensure the analytics provider is GDPR-compliant (e.g., DPF-certified or SCC-backed).

Third-party cookies:

- OmniParse does not set third-party cookies on the `omniparse-ai.vercel.app` domain. No advertising, tracking, social-media, or fingerprinting cookies are used.
- **Note about Stripe:** If you upgrade to a paid plan, Stripe's checkout page (at `checkout.stripe.com`) may set its own cookies on the Stripe domain (not `omniparse-ai.vercel.app`). These are governed by **Stripe's cookie policy** (<https://stripe.com/cookies>) and are necessary for payment security and fraud prevention. OmniParse does not control Stripe's cookies.

Session-only cookies (Next.js framework):

- The Next.js framework may use short-lived session cookies (e.g., for CSRF protection) that are deleted when you close your browser. These are strictly necessary for security and are exempt from consent under ePrivacy Art. 5(3).

3. Legal Basis

Under the EU ePrivacy Directive (2002/58/EC), as amended by Directive 2009/136/EC, and GDPR:

- **Essential storage (op_token, omniparse_cookie_consent, Next.js session cookies):** Exempt from consent requirements as strictly necessary for the provision of the Service explicitly requested by the user (ePrivacy Art. 5(3) exception, as implemented by Czech Act No. 127/2005 Coll. § 89).
- **Non-essential storage (theme, op_shortcuts, op_legal_consent, op_crash_log):** Requires your explicit, informed, prior consent via the cookie banner (GDPR Art. 6(1)(a) + ePrivacy Art. 5(3)). Consent must be as easy to withdraw as to give. You may withdraw consent at any time by (a) clicking the "Essential only" button in the cookie banner re-openable via Settings → Privacy, or (b) clearing your browser's Local Storage.
- You have the right to be informed about what data is collected and for what purpose (this section).
- You have the right to lodge a complaint with the Czech DPA (ÚOOÚ) if you believe our cookie practices violate applicable law.

4. Cookie Banner and Granular Consent

On your first visit to the Service, a cookie banner appears at the bottom of the screen. You have the following options:

- **"Accept all"**: Records consent for both essential and all non-essential storage (theme, shortcuts, legal consent, crash logs).
- **"Essential only"**: Records consent for essential storage only; non-essential items are not set, and any previously set non-essential items are cleared.
- **"Manage preferences" (planned)**: Granular per-category toggles for theme, shortcuts, legal consent, and crash logs. Each category can be accepted or declined independently, in line with EDPB Guidelines 03/2022 on dark patterns.

The banner does not block page load (essential cookies work regardless). Your choice is recorded in Local Storage under the key `omniparse_cookie_consent` along with a timestamp, and persists for 12 months. After 12 months, or if you clear your Local Storage, the banner will re-appear.

You can change your choice at any time by clicking the "Cookie settings" link in the footer, or by clearing the `omniparse_cookie_consent` key from your browser's Local Storage.

5. Managing Your Preferences

You can manage or remove Local Storage data through your browser settings:

- **Chrome**: Settings > Privacy and security > Cookies and other site data > Site data
- **Firefox**: Settings > Privacy & Security > Cookies and Site Data
- **Safari**: Preferences > Privacy > Manage Website Data

Blocking or deleting Local Storage data may affect the Service's functionality — for example, your authentication session and theme preference will be lost.

6. Czech Republic

Under Czech Act No. 127/2005 Coll. (Electronic Communications Act) § 89, as amended effective 1 January 2022 to align with the ePrivacy Directive: the processing of Local Storage data and cookies requires the user's consent unless it is strictly necessary for the operation of the service requested by the user. OmniParse complies with this requirement through the cookie consent banner (§ 4 above). Act No. 480/2004 Coll. (Information Society Services Act) applies to unsolicited commercial communications and is not relevant to OmniParse's cookie practices, as OmniParse does not engage in direct marketing.

7. Contact

For questions about this Cookie Policy: damr58h@gmail.com

Complies with EU ePrivacy Directive 2002/58/EC (as amended by Directive 2009/136/EC), GDPR Art. 6(1)(a) and Art. 7, Czech Act No. 127/2005 Coll. § 89 (Electronic Communications — ePrivacy implementation), and EDPB Guidelines 03/2022 on dark patterns in social media platform interfaces.