

OmniParse

AI-powered invoice extraction & verification

Legal Documents

Transparency, Terms, Privacy & Cookie Policy

VERSION

3.6 (Admin) · Revised October 3, 2026

ISSUED BY

Simon Curda · Natural person · Czech Republic

CONTACT

damr58h@gmail.com

Table of Contents

This document consolidates the four legal instruments that govern your use of OmniParse. The AI Transparency Notice explains how the Service uses artificial intelligence. The Terms of Service set out the binding contract between you and Simon Curda. The Privacy Policy (including the embedded Data Processing Addendum) describes how personal data is handled. The Cookie Policy explains the storage used by the Service.

AI Transparency Notice

Last updated: September 19, 2026 · EU AI Act Article 50 disclosure

Terms of Service

Last updated: September 19, 2026 · Binding contract for use of the Service

Privacy Policy (incl. DPA)

Last updated: September 19, 2026 · GDPR compliance & Data Processing Addendum

Cookie Policy

Last updated: September 19, 2026 · Local storage & consent

Quick summary of v3.5 changes: This version applies the P0 and P1 fixes from the latest legal review. Age requirement raised to 18 for B2B SaaS (was 15), with parental-consent carve-out where applicable law permits. DPA restructured: added §11.1 confidentiality clause, new §11.2 enumerates subprocessor authorization, §11.7 adds subprocessor change-notification mechanism with 30-day notice and objection rights. Retention tightened: personal data deleted within 90 days of account termination. Cookie Policy enumerates each local-storage item with purpose and consent category.

AI Transparency Notice

Last updated: September 19, 2026

EU AI Act Compliance (Regulation EU 2024/1689)

Per Article 50 of the EU AI Act (applicable since **2 August 2026**), deployers of AI systems must inform natural persons that they are interacting with an artificial intelligence system. This notice fulfills that obligation. Article 5 prohibitions on certain AI practices have been applicable since **2 February 2025**; OmniParse confirms that none of its processing falls within any Article 5 prohibited category (social scoring, manipulative AI, untargeted facial scraping, emotion recognition in workplaces/schools, etc.).

1. AI Disclosure

OmniParse uses artificial intelligence. Specifically:

- **Document parsing:** A vision-language model (VLM) analyzes uploaded documents (PDFs, images) to extract structured invoice data.
- **Chat assistant:** A large language model (LLM) generates responses to your questions about your invoice data.
- **Artifact generation:** The chat AI may generate structured data (tables, charts, summaries) based on your requests.

You are interacting with AI-generated output throughout the Service. All parsing results, chat responses, and generated artifacts are produced by AI models and should be treated accordingly.

2. System Classification

Based on the current intended purpose and functionality of the Service, OmniParse does not currently identify a use case that falls within the high-risk categories listed in Annex III of the EU AI Act.

The current Service is intended for invoice extraction, document analysis and conversational assistance. It is not intended to perform biometric identification, employment or education decisions, eligibility decisions for essential public or private services, law-enforcement assessments, migration or border-control decisions, or decisions in the administration of justice or democratic processes.

The classification of any future feature depends on its specific intended purpose, functionality and manner of deployment and will be assessed separately where required.

The Service is subject to applicable transparency requirements under Article 50 of the EU AI Act, in particular Article 50(1) (AI systems that interact with natural persons). The Service may also generate or manipulate content using AI. Where applicable, we will implement the transparency measures required by Article 50 for AI-generated or manipulated content (Article 50(4)). We do not assert that every output of the Service automatically falls within all Article 50(4) sub-categories; the applicability of specific transparency obligations is assessed based on the nature of the output and the context of its use.

As the provider of the Service under our own name and brand, we acknowledge our responsibilities under the EU AI Act. This assessment may change if the Service is used in high-risk contexts. We will update this notice accordingly and implement the required obligations if reclassification occurs.

General-Purpose AI (GPAI) classification: The AI models used by OmniParse (e.g., Gemini, Pixtral, Llama) are developed by upstream providers (Google, Mistral, Meta, etc.) and may qualify as GPAI models under Art. 3(63). OmniParse's AI Act role depends on the specific component and context: OmniParse may act as a **deployer** of these models (within the meaning of Art. 3(4)) when it uses them under its own name for the invoice-extraction and chat-assistant purposes. For some components — particularly where OmniParse integrates the model into a broader service offered under its own brand and substantially shapes the system's intended purpose — OmniParse's role may also be that of a **downstream provider** of an AI system within the meaning of Art. 25. The exact classification of each component is fact-specific and should be confirmed by legal counsel. Transparency obligations applicable to GPAI providers under Art. 53 rest with the upstream model providers (Google, Mistral, Meta, etc.). OmniParse relies on the technical documentation and information provided by those upstream providers.

3. Technology Details

- **Type:** Document understanding and natural language processing
- **Components:** Vision-language model (VLM) for document parsing, Language model (LLM) for chat
- **Purpose:** Extract structured data from documents and provide conversational analysis
- **Data processed:** Invoices and related documents may contain personal data (e.g. names, email addresses, phone numbers, bank account details, signatures). Depending on how the Service is used, we may process personal data on behalf of users or their organizations.
- **AI providers:** We may use multiple third-party AI providers to process uploaded content. The provider used for a particular request may depend on availability, capacity, configuration, and other operational factors. A current list of subprocessors and applicable processing locations is available in our Privacy Policy.

4. Known Limitations

The AI may make mistakes. Specific risk areas include:

- **Inaccurate extraction:** Low-quality scans, unusual layouts, handwritten text, or non-English documents may produce errors.
- **Hallucinations in chat:** The AI may generate plausible but factually incorrect information about your data.
- **Artifact errors:** Generated tables, charts, and summaries inherit any extraction or reasoning errors.
- **Currency and formatting:** May misinterpret currency symbols, date formats, or number formatting across regions.
- **Calculation errors:** Calculations and statistics generated by the Service may contain errors, including errors caused by incorrect extraction, interpretation, aggregation, rounding, currency conversion or other processing. You are responsible for independently verifying results before relying on them.
- **Confidence indicators:** AI-generated confidence indicators are heuristic signals intended to help prioritize manual review. They are not statistical probabilities and do not constitute a guarantee of accuracy.
- **Provider availability:** AI providers may impose rate limits, discontinue models, or become unavailable. During periods of high demand, the Service may return "temporarily busy" messages.

Always verify AI-generated data against your original documents before using it for any purpose where accuracy matters (financial reporting, tax filing, legal proceedings).

5. No Professional Advice

The Service is provided for informational and administrative purposes only and does not constitute accounting, tax, legal, financial or other professional advice.

AI-generated calculations, classifications, summaries, recommendations and chat responses must not be relied upon as the sole basis for financial, accounting, tax, legal or business decisions. Always consult a qualified professional for advice specific to your situation.

Currency conversions, exchange rates, tax calculations and financial aggregations may be inaccurate or based on incomplete or outdated information.

6. Human Oversight

OmniParse provides the following human oversight mechanisms (consistent with AI Act Art. 14 best practices, even though the Service is not high-risk):

- **Confidence indicators:** Extractions include heuristic confidence indicators to help prioritize manual review. These are not statistical probabilities.
- **Manual review:** Users can review, edit, or delete any extracted data at any time.
- **No autonomous high-impact decisions:** In the current production configuration, AI features do not independently approve payments, execute payments, or make decisions about individuals with legal or similarly significant effects. AI extraction, classification and chat outputs are presented as assistive results for user review. Users remain responsible for reviewing and approving material business actions.
- **Future functionality:** If OmniParse introduces features that independently make or materially influence decisions with legal or similarly significant effects concerning individuals, those features will be separately assessed under applicable GDPR and EU AI Act requirements before deployment.
- **Feedback loop:** Users can re-upload documents, correct results, and delete inaccurate data.
- **Audit trail:** All user actions (edits, approvals, deletions) are logged in the user's account audit log for traceability.

7. Data Handling and AI Training

- **Processing by AI providers:** Documents and chat messages are sent to AI providers for inference. Data may be temporarily processed, cached, or retained by AI providers for security, abuse prevention, monitoring, billing, debugging, or other purposes specified in their applicable terms and data-processing agreements. Retention periods and processing locations vary by provider and service configuration.
- **Stored on our side:** We retain uploaded documents for up to 30 days in the active application environment. Backup copies, security logs and data processed by third-party providers may be retained for different periods where necessary for security, legal or operational purposes, as described in our Privacy Policy.
- **Model training — production configuration:** OmniParse does not intentionally route customer content through AI API configurations that are documented by the provider as permitting provider training on customer content when the default production configuration is in effect. Specifically: (a) OmniParse sends `usage_options={"enable_training": false}` on every Mistral API call when `MISTRAL_DISABLE_TRAINING=true` (the default). This setting is effective on the applicable paid Mistral tier; on a free tier, Mistral may ignore or reject the parameter, and the parameter itself is not represented as an independent guarantee of provider-level data handling; (b) OpenRouter is disabled by

default and is not used for customer content unless the operator explicitly enables it after completing the applicable contractual, data-protection and transfer review; (c) Groq processing is governed by the applicable Groq contractual and data-protection terms; and (d) Google Gemini through the Google AI Studio free-tier endpoint is disabled by default and is not used unless the operator explicitly enables it after reviewing the applicable Google AI Studio terms. Operators who change the default provider configuration are responsible for verifying the applicable provider terms, transfer mechanism, retention and training settings and, where necessary, updating the Privacy Policy and this AI Transparency Notice before processing customer content through the changed configuration.

- **Important note on free-tier AI APIs:** Some AI providers may have different data handling terms for free/unpaid API tiers compared to paid tiers. We recommend reviewing the applicable provider terms for details. For production use involving personal data, paid API tiers may provide stronger data protection guarantees.
- **Metadata logging:** We log API request metadata (timestamps, success/failure, provider used) for operational purposes. We do not log the content of your documents to server logs.

8. Model and Provider Changes

AI models, providers and processing methods may be changed, updated, replaced or discontinued. Material changes to AI providers (e.g., switching the primary vision-extraction provider) will be announced in the Service or by email at least 7 days in advance, except where the change is urgently required for security or availability reasons. Changes in models may affect extraction results, classifications, calculations and chat responses. The same input document may produce different results when processed by different models or at different times.

The Service is provided on an "as is" and "as available" basis. We do not guarantee uninterrupted, error-free or continuously available operation of the Service or any AI provider. AI models and third-party providers may change, become unavailable, impose rate limits, discontinue models or modify their capabilities. Material changes will be notified per the above; immaterial changes (e.g., model version bumps within the same family) may happen without separate notice but will be reflected in the "Last updated" date of this Notice.

9. Your Rights as an AI System User

- Review original documents alongside extracted data
- Manually correct any AI-extracted fields
- Delete and re-upload documents to get new extractions
- Request information about the processing and factors relevant to an AI-generated result, where technically and legally available (contact us)
- Lodge a complaint about AI output quality (damr58h@gmail.com)
- You may stop using AI-powered processing and delete your data. Where supported, you may export your data for processing outside the Service.

10. International Transfers

AI processing in the default production configuration is performed by Mistral AI (EU-based, Paris, France) and Groq Inc. (US-based).

OpenRouter and Google Gemini through Google AI Studio are disabled by default and are not used unless the operator explicitly enables them after completing the applicable contractual, data-protection and transfer review.

Where personal data is transferred to a third country, OmniParse relies on an applicable lawful transfer mechanism under GDPR Chapter V, such as an adequacy decision, Standard Contractual Clauses (SCCs), or another lawful mechanism where applicable.

For US recipients, the EU-US Data Privacy Framework (DPF) may provide an applicable transfer mechanism where the recipient is covered by the relevant DPF certification. Where an adequacy mechanism does not apply, OmniParse relies on an appropriate GDPR Chapter V safeguard, such as SCCs, together with supplementary measures or a Transfer Impact Assessment where appropriate.

Google AI Studio free-tier processing is subject to Google's applicable AI Studio terms. The Google Cloud Data Processing Addendum applies to the applicable Google Cloud / Vertex AI services and should not be interpreted as automatically applying to Google AI Studio free-tier processing. This is one reason Google Gemini is disabled by default in the production configuration.

The primary database is hosted by Supabase in Ireland (EU). The applicable processing locations and transfer safeguards are described in the Privacy Policy.

11a. AI Literacy (AI Act Art. 4)

Pursuant to Article 4 of the EU AI Act (applicable since 2 February 2025), providers and deployers of AI systems shall take measures to ensure, to their best extent, a sufficient level of AI literacy in their staff and persons dealing with the operation and use of AI systems on their behalf. This Notice forms part of OmniParse's measures supporting AI literacy among persons involved in the operation and use of the Service. It is not, by itself, a complete fulfillment of the Article 4 obligation — that obligation also encompasses internal staff training, operational documentation, and ongoing assessment of AI literacy needs. Internal staff training on AI capabilities, limitations, and risks is documented in our internal records and updated at least annually, or when significant changes to AI features occur.

11b. Logging and Monitoring

Although OmniParse is not classified as a high-risk AI system (and therefore Article 12 logging requirements do not strictly apply), we maintain the following logging practices as best practice:

- **Application-level audit log:** User actions (invoice create/edit/approve/delete, chat send) are logged in the user's account and retained for the lifetime of the account.
- **Server-side operational logs:** Vercel function logs retained for 30 days for debugging and incident investigation.
- **AI request metadata:** Timestamps, provider used, success/failure status, and request duration are logged for operational monitoring. The **content** of documents and chat messages is not logged to server logs.
- **Incident logs:** Security incidents involving AI systems are documented and retained for at least 2 years for regulatory audit purposes.

12. Contact

- **AI ethics concerns:** damr58h@gmail.com (subject line "AI Ethics")

- **Data protection / GDPR:** damr58h@gmail.com (subject line "DSR" or "DPA")
- **Legal:** damr58h@gmail.com
- **Czech DPA (ÚOOÚ):** uoou.cz (<https://www.uoou.cz>)

The Service is intended to comply with applicable laws in the jurisdictions in which it is offered. Users remain responsible for ensuring that their use of the Service complies with applicable local requirements.

Terms of Service

Last updated: September 19, 2026

1. Acceptance of Terms

By accessing or using OmniParse ("the Service"), you agree to be bound by these Terms of Service ("Terms"). If you do not agree with any part of these Terms, you must not use the Service. These Terms apply to all users of the Service.

When you register an account, you are required to explicitly accept these Terms by checking the "I agree to the Terms of Service and Privacy Policy" checkbox. The timestamp of your acceptance is recorded in your account record. Continued use of the Service after any changes to these Terms constitutes acceptance of the updated Terms.

2. Description of Service

OmniParse is an AI-powered document parsing application that extracts structured data from invoices and similar business documents. The Service uses vision-language models to analyze uploaded documents and return extracted data in JSON format.

Current capabilities:

- Document upload (PDF, JPEG, PNG, WebP) up to 20MB per file
- AI-powered field extraction with per-field confidence scores
- AI chat assistant with interactive artifact generation (tables, charts, summaries)
- Data export (CSV and JSON) from the invoices table
- Visual analytics dashboard with charts
- User authentication with email, password, and JWT tokens
- GDPR-compliant data export and account deletion

3. Eligibility

The Service is intended for business and professional users and is not directed at children. You must be at least **18 years old** to register an account, unless applicable law permits use by a younger person with the required parental or guardian consent and the contractual arrangement is legally valid. In the Czech Republic, the age of digital consent under GDPR Art. 8 (as implemented by Czech Act No. 110/2019 Coll.) is 15; in other EU/EEA Member States the age ranges from 13 to 16 depending on national law.

Where applicable law permits use by a person aged between the applicable digital-consent age and 18, that person must obtain parental or guardian consent to use this Service and to enter into these Terms. Where a minor accesses the Service through an organizational account (e.g., as an employee or contractor of a business customer), the organization is responsible for ensuring that such access is lawful and complies with applicable child-labor and data-protection law.

If you are registering on behalf of a business entity, you confirm that you have the authority to bind that entity to these Terms.

4. User Responsibilities and Warranties

You represent and warrant that you have all rights, permissions and lawful authority necessary to upload, process and transmit documents, emails and other content through the Service.

- You are solely responsible for the documents you upload to the Service.
- You must have the legal right to process any documents you submit (ownership, authorization, or legitimate interest).
- You must not upload documents containing personal data of EU data subjects unless you have a legal basis to do so under GDPR.
- Unless explicitly supported by the Service, users must not upload documents containing special categories of personal data (such as health information, biometric data, information revealing religious or political beliefs, or other highly sensitive information).
- When using email import, the Service may process email content, headers, attachments and other information necessary to identify and process invoices. Users should not forward emails containing information that they are not authorized to disclose to the Service.
- You must not use the Service for any unlawful purpose, including fraud, money laundering, or tax evasion.
- You must not attempt to reverse-engineer, decompile, or disrupt the Service.
- You must not use automated tools to overwhelm the Service (no scraping, no DDoS, no rate-limit evasion).
- You must not attempt to extract other users' data through IDOR attacks or similar exploits.
- You acknowledge that AI-extracted data may contain errors and you are responsible for verifying results before use.
- You are responsible for keeping your account password confidential.
- You are responsible for maintaining appropriate backups of documents and data that are important to you. The Service should not be considered the sole repository or archival system for financial records.

5. Intellectual Property and License

Your content: Content you upload (documents, chat messages, extracted data) remains your property. OmniParse does not claim ownership of your documents or extracted data. By uploading content, you grant OmniParse a worldwide, non-exclusive, royalty-free, revocable license to process, transmit, store, and display your content solely for the purpose of operating and providing the Service to you, including the right to transmit content to third-party AI inference providers as described in the Privacy Policy. This license ends when the relevant content is deleted from the Service, except to the extent continued processing or retention is required for legitimate operational, security, legal, or backup purposes described in the Privacy Policy (including backup copies that may persist for a limited period after deletion from active storage).

Service IP: The Service itself, including its source code, design, UI, documentation, brand, and underlying infrastructure, is the property of Simon Curda. You may not copy, modify, distribute, or create derivative works from the Service itself without written permission. The names "OmniParse" and associated logos are trademarks of Simon Curda.

Feedback: If you provide feedback, suggestions, or ideas about the Service, you grant OmniParse a perpetual, irrevocable, royalty-free license to use and incorporate that feedback into the Service without obligation or compensation to you.

6. Data Processing

Uploaded documents are retained in the active application environment for up to 30 days for preview purposes and are then deleted from active storage, subject to applicable backup, security, legal and provider retention periods described in our **Privacy Policy** (/privacy-policy) (Privacy Policy §4 retention table). Extracted data is stored in your account database and is accessible from your dashboard. You can delete your data at any time. See our **Privacy Policy** (/privacy-policy) for full details on data handling, including the distinction between active-storage deletion (30 days) and post-termination cleanup (90 days).

7. AI Disclaimer

THE SERVICE USES ARTIFICIAL INTELLIGENCE FOR DOCUMENT ANALYSIS, DATA EXTRACTION, AND CONVERSATIONAL CHAT RESPONSES. AI OUTPUT MAY CONTAIN ERRORS, INACCURACIES, OR HALLUCINATIONS.

You must always verify AI-extracted data against the original documents before using it for financial reporting, tax filing, legal proceedings, or any context where accuracy is critical. OmniParse provides confidence scores as a guide, but these are estimates and not guarantees of accuracy.

The AI Chat assistant is a conversational tool, not a source of truth. Responses to your questions about your invoice data are generated by AI models and may be inaccurate, incomplete, or out-of-context. Always cross-reference chat responses against the underlying invoice records and source documents before making financial decisions, payments, vendor communications, or tax filings based on them.

AI providers: OmniParse uses third-party AI providers for vision extraction and text generation. The primary provider is **Mistral AI** (Paris, France — EU-based; transfers are expected to remain within the EEA, subject to Mistral's applicable terms). **Groq Inc.** (US-based) has confirmed Standard Contractual Clauses (SCCs) in effect (EU SCC Module 2, self-executing under Groq's DPA dated October 15, 2025). **OpenRouter** (US-based) and **Google Gemini** (US-based, AI Studio free tier) are **disabled by default** and are not used unless the operator explicitly enables them after completing their own DPA/SCC review (OpenRouter) or AI Studio terms review (Google Gemini). These providers may experience rate limits, outages, or model deprecations that affect the availability and quality of AI features. We do not control and are not responsible for the behavior of these upstream models beyond our integration layer.

See our **AI Transparency Notice** (/ai-act-notice) for full details on AI limitations and your rights.

7a. No Professional Advice

The Service is provided for informational and administrative purposes only and does not constitute accounting, tax, legal, financial or other professional advice. AI-generated or AI-extracted information may contain errors, omissions or hallucinations. Users must verify extracted amounts, dates, tax information, bank details and other material fields before relying on them. AI-generated calculations, classifications, summaries, recommendations and chat responses must not be relied upon as the sole basis for financial, accounting, tax, legal or business decisions.

Currency conversions, exchange rates, tax calculations and financial aggregations may be inaccurate or based on incomplete or outdated information. Always consult a qualified professional for advice specific to your situation. OmniParse is an assistive processing tool and does not provide accounting, tax, legal or financial advice. Users remain responsible for reviewing outputs and making their own business decisions.

Where the Service provides features that automate invoice approval, categorization, or other workflow decisions, such features are configurable and require user setup and review. If any automated feature produces a decision with legal or similarly significant effects concerning an individual (e.g., a vendor or employee), the applicability of GDPR Art. 22 and the EU AI Act must be assessed for that specific function. OmniParse does not represent that a single disclaimer is sufficient for all possible automated-decision scenarios; each material feature must be assessed on its own facts.

7b. Calculation and Data Accuracy Disclaimer

Calculations and statistics generated by the Service may contain errors, including errors caused by incorrect extraction, interpretation, aggregation, rounding, currency conversion or other processing. You are responsible for independently verifying all results before relying on them for any purpose.

AI models, providers and processing methods may be changed, updated, replaced or discontinued. Material changes to AI providers (e.g., switching the primary vision-extraction provider) will be announced in the Service or by email at least 7 days in advance, except where the change is urgently required for security or availability reasons. Changes in models may affect extraction results, classifications, calculations and chat responses. The same input document may produce different results when processed by different models or at different times.

7c. No Payment Authorization

OmniParse is an invoice intelligence platform, not a payment system. Unless explicitly enabled by you through a supported integration (which OmniParse does not currently offer), OmniParse does not initiate, authorize, execute, or otherwise cause any payment, bank transfer, wire transfer, ACH transaction, credit-card charge, or other financial transaction. The Service extracts, organizes, and analyzes invoice data — it does not pay invoices.

Any payment you make to a vendor, supplier, or other recipient — even if OmniParse extracted the vendor's bank details, suggested a payment amount, or displayed a "Pay" button that links to an external site — is solely your responsibility. You must independently verify the payee, amount, bank account, and payment method before initiating any payment. To the maximum extent permitted by applicable law, OmniParse is not a party to any payment transaction and bears no liability for payments made, missed, or incorrectly executed based on data extracted or suggested by the Service. Nothing in this section excludes or limits liability that cannot lawfully be excluded or limited under applicable law, including liability for intentional damage, gross

negligence, and mandatory consumer protection rights (see §10.4 and §10.5).

If OmniParse adds payment-integration features in the future (e.g., bank integrations, accounting-system integrations, payment links), such features will be governed by separate terms specific to those integrations and will require your explicit opt-in. This section applies to the current Service, which does not include payment-execution capabilities.

8. Service Availability and Limitations (Including IMAP Email Capture)

The Service is provided on a best-effort basis. OmniParse uses automated processes (including IMAP email scanning, AI vision-language models, classification heuristics, and external AI providers) to capture and extract invoice data. None of these processes are guaranteed to be complete, accurate, or timely. We explicitly disclaim any warranty that the Service will capture every invoice, classify it correctly, or extract all data without error.

The IMAP email auto-capture feature is an auxiliary convenience, not a substitute for proper invoice management. Specifically, the Service may fail to capture invoices for reasons including but not limited to:

- AI provider rate limits, outages, model deprecations, or temporary unavailability
- Classifier false negatives (the AI incorrectly determines an email is not an invoice)
- Email attachment format issues (corrupted PDFs, unsupported image formats, encrypted archives)
- IMAP server-side issues (mailbox sync delays, throttling by Gmail/Outlook/etc.)
- Vercel function timeout (60-second limit per scan call)
- Network connectivity issues between Vercel, the email provider, and AI providers
- Pending Review queue cap (100 items per inbox — scan pauses when reached)
- Sender blocklist (user or system blocks a legitimate invoice sender)
- Emails older than the scan cursor (UID-based resume may skip emails if cursor is corrupted)
- Bugs, configuration errors, or schema migrations that affect data integrity

You must not rely solely on the Service for time-critical or financially significant invoice management. You remain responsible for:

- Periodically checking that all expected invoices have been captured
- Manually uploading invoices that the Service failed to capture
- Verifying due dates and payment deadlines against source documents
- Maintaining backup processes for invoice tracking independent of OmniParse
- Reviewing the Pending Review queue regularly (at minimum weekly, or before any payment run)

To the maximum extent permitted by applicable law, OmniParse shall not be liable for losses resulting from missed, delayed, incorrectly classified, or incorrectly extracted invoices where such losses arise from limitations, errors, or failures of the Service. This includes late payment fees, vendor relationship damage, tax filing errors, late filing penalties, or any other damages arising from the Service's failure to capture or correctly process an invoice, including scenarios where the user has stopped checking their email inbox manually and relies exclusively on the Service's auto-capture feature. Nothing in this section excludes or limits liability that cannot lawfully be excluded or limited under applicable law, including liability for intentional damage, gross negligence, and mandatory consumer protection rights (see §10.4 and §10.5).

For paid users only: Where OmniParse is found directly at fault for damages (e.g., a confirmed bug in our code, not a third-party AI provider issue), liability is capped at the amount you paid to OmniParse in the 12

months preceding the claim. For Free tier users, liability for slight negligence is excluded to the extent permitted by applicable law; liability for intentional damage and gross negligence remains unaffected in all cases. Nothing in this section excludes or limits liability that cannot lawfully be excluded or limited under applicable law, including mandatory consumer-protection rights under Czech Act No. 634/1992 Coll. and EU Directive 93/13/EEC.

Notification of issues: If you become aware that the Service has failed to capture an invoice or has extracted incorrect data, please notify us at **damr58h@gmail.com** as soon as practicable. We will investigate and, where possible, restore or correct the affected data. Prompt notification helps us serve you better; this does not shorten any statutory limitation period under applicable law.

9. Disclaimer of Warranties

THE SERVICE IS PROVIDED "AS IS" AND "AS AVAILABLE" WITHOUT WARRANTIES OF ANY KIND, EITHER EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO IMPLIED WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE, AND NON-INFRINGEMENT.

We do not warrant that the Service will be uninterrupted, error-free, or free of harmful components. We do not warrant that AI extraction results will be accurate or complete. AI providers, including Mistral and Groq, may experience rate limits, model deprecations, or outages that affect availability of AI features.

Mistral AI is based in Paris, France. Groq Inc. is based in the United States. OpenRouter and Google Gemini are **disabled by default** and are not used for customer content unless explicitly enabled by the operator following the applicable contractual, data-protection and transfer review.

10. Limitation of Liability

This Section 10 sets out the limitation of liability framework between you and Simon Curda (the operator of OmniParse). It is structured as: (a) direct damages, (b) indirect/consequential damages, (c) aggregate liability cap, (d) exceptions, (e) mandatory statutory rights.

10.1 Direct Damages

To the maximum extent permitted by applicable law, Simon Curda's liability for direct damages arising from your use of the Service is capped at the amount you paid to OmniParse in the 12 months preceding the claim. For Free tier users (who paid nothing), liability for direct damages is capped at zero to the extent permitted by applicable law.

10.2 Indirect and Consequential Damages

To the maximum extent permitted by applicable law, Simon Curda shall not be liable for any indirect, incidental, special, consequential, or punitive damages, including but not limited to loss of profits, loss of data, loss of business opportunities, reputational damage, or costs of substitute services. This includes losses arising from reliance on AI-generated output that you failed to independently verify against source documents.

10.3 Aggregate Liability Cap

The aggregate liability of Simon Curda for all claims arising out of or relating to these Terms or the Service (whether in contract, tort, or otherwise) shall not exceed the total amount you paid to OmniParse in the 12 months preceding the claim. For Free tier users, the aggregate liability cap is zero to the extent permitted by applicable law.

10.4 Exceptions (Liability Not Limited)

The limitations in §10.1, §10.2, and §10.3 do **not** apply to liability that cannot lawfully be excluded or limited under applicable law, including:

- Liability for intentional damage or intentional misconduct;
- Liability for gross negligence under applicable mandatory provisions of Czech law;
- Liability for personal injury or death caused by negligence;
- Liability arising from breach of statutory duties under mandatory consumer-protection law;
- Liability for defects in goods or services to the extent such liability cannot be contractually excluded under applicable mandatory law;
- Liability under product liability laws;
- Any other liability that cannot be excluded or limited under applicable mandatory law.

10.5 Mandatory Consumer Rights (EU / Czech)

EU consumer note: If you are a consumer within the meaning of Czech Act No. 89/2012 Coll. (Civil Code) § 419, nothing in these Terms limits any statutory consumer rights you may have under EU or Czech law, including under:

- Czech Act No. 89/2012 Coll. (Civil Code) — contractual liability, defects, consumer protection;
- Czech Act No. 634/1992 Coll. (Consumer Protection Act) — out-of-court dispute resolution;
- EU Directive 93/13/EEC (Unfair Contract Terms) — unfair terms are not binding on consumers;
- EU Directive 2011/83/EU (Consumer Rights) — right of withdrawal, pre-contractual information;
- GDPR (EU Regulation 2016/679) — data subject rights, liability under Art. 82.

Where mandatory consumer-protection law applies, it overrides conflicting provisions of these Terms to the extent of the conflict. If any provision of §10.1–10.3 is found unenforceable against a consumer, that provision shall be limited or eliminated to the minimum extent necessary to comply with applicable law, and the remaining provisions shall remain in full force and effect.

10.6 Basis of Limitation

The limitations in this Section 10 reflect the allocation of risk between the parties and form an essential basis of the bargain. The Service is offered at a price that reflects these limitations, and you acknowledge that the price would be substantially higher absent these limitations. Operators who require higher liability caps (e.g., Enterprise customers with negotiated contracts) may contact damr58h@gmail.com to discuss custom terms.

11. Indemnification

To the extent permitted by applicable law, and except where prohibited by mandatory consumer protection law (including EU Directive 93/13/EEC on unfair contract terms and Czech Act No. 634/1992 Coll.), you agree to indemnify and hold harmless Simon Curda from and against any claims, damages, losses, or expenses (including reasonable attorneys' fees) arising from (a) your use of the Service in violation of these Terms or applicable law, (b) your infringement of any third-party rights (including uploading documents you do not have the right to process), or (c) your violation of applicable data-protection law in respect of personal data you upload.

Consumer carve-out: If you are a consumer within the meaning of Czech Act No. 89/2012 Coll. (Civil Code) § 419, this indemnification clause applies only to the extent that the relevant claim arises from your intentional

breach or gross negligence, and in no event shall a consumer be required to indemnify Simon Curda for damages caused by Simon Curda's own act or omission.

12. Subscription Plans and Billing

Paid plans (Pro, Plus, Business, Enterprise) are billed monthly or annually via Stripe. You can cancel your subscription at any time from your Stripe customer portal. Cancellation takes effect at the end of the current billing period — you retain access to paid features until then, and no further charges are made.

When you delete your account, any active Stripe subscription is automatically cancelled to prevent ongoing charges.

Refund policy:

- **Within 14 days of payment, no AI usage:** Full refund on request (see §14 Right of Withdrawal). Contact damr58h@gmail.com.
- Within 14 days of payment: Consumers retain any statutory right of withdrawal unless and until the applicable statutory conditions for exclusion or loss of that right have been satisfied.
- Where the consumer has expressly requested that performance begin during the withdrawal period and has provided the legally required acknowledgment concerning the consequences of doing so, the right of withdrawal may cease or be limited to the extent permitted by applicable law.
- Nothing in this refund policy limits mandatory statutory consumer rights.
- **After 14 days, monthly plan:** No refund of the current month's charge; subscription remains active until end of billing period. No further charges after cancellation.
- **After 14 days, annual plan:** Pro-rata refund of unused full months remaining on the plan, less a reasonable administrative fee of EUR 10. Contact damr58h@gmail.com to request a pro-rata refund.
- **Service-impacting outage:** If OmniParse experiences a verified outage exceeding 24 consecutive hours that prevents you from using paid features, contact us for a service-credit or pro-rata refund for the affected period.
- **Free tier:** No payment is made, so no refund is applicable. Account deletion is available at any time.

Refunds are processed back to the original payment method via Stripe within 14 days of approval, as required by Art. 13(1) of Directive 2011/83/EU.

13. Modifications to Terms

We may update these Terms from time to time. **Material changes will be notified to you via email (to your account email address) and via an in-Service notice at least 30 days before the change takes effect.** For changes that are not material (e.g., corrections of typos, clarifications, or changes that do not affect your rights or obligations), we may update these Terms without prior notice, with effect from the "Last updated" date.

For consumer users: Material changes to these Terms that adversely affect your rights or obligations will not be treated as expressly accepted solely because you continue using the Service. Where applicable law requires express agreement to a material change, OmniParse will obtain that agreement before the change applies to you.

Where applicable law permits changes to continuing contracts subject to notice and a right to terminate, OmniParse will provide the required notice and termination opportunity. If you do not agree to a material change, you may terminate your account before the change takes effect, subject to any applicable statutory rights.

The "Last updated" date at the top of this page indicates when these Terms were last revised. We will retain prior versions of these Terms and provide them on request.

14. Right of Withdrawal (EU Consumers)

Under Czech Civil Code (Act No. 89/2012 Coll.) §§ 1829 and 1837, and Directive 2011/83/EU on consumer rights, consumers concluding distance contracts (contracts concluded online) have a **14-day right of withdrawal**.

However, pursuant to § 1837(j) of the Civil Code and Art. 16(m) of Directive 2011/83/EU, the right of withdrawal does **NOT** apply to:

- Digital content or services where the consumer has expressly requested that the service begin during the withdrawal period AND has acknowledged that by doing so, they lose the right to withdraw;
- Services that have been fully performed with the consumer's prior express consent and acknowledgment.

For this waiver to be valid under § 1837(j) and Art. 16(m) CRD, the following express acknowledgment is required at signup (separate from the Terms acceptance checkbox):

"I acknowledge that OmniParse is a digital service that begins immediately upon my first use of AI features (upload, scan, or chat). I expressly consent to OmniParse beginning performance during the 14-day withdrawal period, and I acknowledge that I will lose my right of withdrawal once the service has started, pursuant to § 1837(j) of the Czech Civil Code and Art. 16(m) of Directive 2011/83/EU."

By creating an OmniParse account and checking this separate acknowledgment box, you expressly agree that:

- The Service may begin immediately upon your first use of AI features (upload, scan, or chat).
- Where the statutory conditions for exclusion of the right of withdrawal are satisfied, you expressly request that performance begin during the withdrawal period and acknowledge the applicable legal consequences described in §1837(j) of the Czech Civil Code and Article 16(m) of Directive 2011/83/EU.
- This acknowledgment does not affect any statutory rights that cannot lawfully be waived.

One-click withdrawal button: Pursuant to Art. 4(2) of Directive (EU) 2023/2673 (effective 19 June 2026), if you have not yet used any AI features, a one-click "Withdraw" button is available in your account Settings that allows you to exercise your right of withdrawal without further formalities. Withdrawal must be confirmed by us within 24 hours.

For Free tier users: Since no payment is made, no withdrawal refund is applicable. Account deletion is available at any time.

For paid subscribers: If you cancel your subscription within 14 days of payment AND have not used any AI features in that period, you are entitled to a full refund (processed within 14 days via Stripe). Contact **damr58h@gmail.com** to request a refund. If you have used AI features during the 14-day period, the right of withdrawal is forfeited per § 1837(j).

15. Consumer Information and Dispute Resolution

Pursuant to § 1753 of the Civil Code (Act No. 89/2012 Coll.), the following information is provided to consumers before concluding a distance contract:

- **Provider:** Simon Curda (natural person), Czech Republic
- **Contact:** damr58h@gmail.com

- **Service description:** AI-powered invoice extraction, validation, and analytics platform
- **Total price (incl. taxes):** Free tier available (no charge). Paid plans from USD 49/month to USD 499/month as described on the pricing page; VAT is reverse-charged for EU B2B customers with a valid VAT ID and included for EU consumers at the applicable Czech VAT rate.
- **Additional fees:** None beyond the subscription price. No delivery costs (digital service delivered online).
- **Contract duration:** Monthly or annual subscription, cancelable at any time (cancellation takes effect at end of billing period).
- **Performance timeframe:** The Service begins immediately upon account creation and first use of AI features.
- **Payment method:** Credit/debit card via Stripe. Invoicing available for Enterprise plans on request.
- **Right of withdrawal:** See § 14 above (limited for digital content under § 1837(j) Civil Code).
- **Applicable law:** Czech law (see § 16 below)
- **Court of jurisdiction:** For consumer disputes, the court of the consumer's domicile; for B2B disputes, the competent courts of the Czech Republic.
- **Language:** These Terms are available in English. Communication with the provider may be conducted in English or Czech.
- **After-sales obligations:** None beyond the warranties and obligations set out in these Terms.
- **Professional liability insurance:** No professional liability insurance is currently maintained. This does not exclude or limit any liability that cannot lawfully be excluded or limited under applicable law.
- **Code of conduct:** The provider has not subscribed to any code of conduct under § 1753(1)(i) Civil Code.

Out-of-Court Dispute Resolution (ADR)

Czech consumers may seek out-of-court dispute resolution through the following entity:

- **Česká obchodní inspekce (ČOI)** — Czech Trade Inspection Authority
- Address: Štěpánská 567, 120 00 Praha 2, Czech Republic
- Website: **coi.gov.cz** (<https://coi.gov.cz>)
- ADR information: **coi.gov.cz/informace-o-adr** (<https://coi.gov.cz/informace-o-adr>)

A proposal for out-of-court dispute resolution must be submitted within 1 year from the date the consumer first exercised their right in the matter. The procedure is governed by Act No. 634/1992 Coll. (Consumer Protection Act) §§ 20b–20e.

EU consumers may also seek alternative dispute resolution through the European Commission's consumer redress resources at **consumer-redress.ec.europa.eu** (<https://consumer-redress.ec.europa.eu>).

16. Governing Law and Dispute Resolution

These Terms are governed by the laws of the Czech Republic, without regard to conflict of law principles. EU consumer protection laws apply additionally for users classified as consumers under EU law.

For disputes arising from these Terms, you agree to first attempt resolution by contacting us at **damr58h@gmail.com**. If the dispute cannot be resolved within 30 days, it shall be submitted to the competent courts of the Czech Republic. For consumer disputes, the court of the consumer's domicile shall have jurisdiction.

17. Severability

If any provision of these Terms is found to be unenforceable or invalid, that provision shall be limited or eliminated to the minimum extent necessary, and the remaining provisions shall remain in full force and effect.

18. Non-Reliance on AI Output

You acknowledge and agree that:

- (a) AI-extracted data and chat responses are provided for informational and administrative purposes only;
- (b) OmniParse does not warrant that AI output will be accurate, complete, timely or suitable for any particular purpose;
- (c) you are responsible for reviewing AI output against the underlying source documents before relying on it;
- (d) confidence scores are model-generated estimates intended to assist with prioritising review and are not guarantees of accuracy;
- (e) AI chat responses may contain errors, hallucinations or information that is incomplete or out of context; and
- (f) OmniParse may modify, suspend or discontinue AI features or change AI providers, subject to the notice requirements described in these Terms.

Any limitation of liability arising from reliance on AI output is subject to §10 and to any liability that cannot lawfully be excluded or limited.

20. Force Majeure

Neither party shall be liable for any failure or delay in performing its obligations under these Terms (except for payment obligations) to the extent such failure or delay is caused by events beyond its reasonable control, including but not limited to acts of God, natural disasters, war, terrorism, civil unrest, epidemics or pandemics, governmental actions, labor disputes, internet or telecommunications failures, or failures of third-party AI providers, hosting providers, or payment processors. The affected party shall notify the other party without undue delay and shall use commercially reasonable efforts to resume performance as soon as reasonably practicable.

21. Notices

Notices from OmniParse to you: Will be sent to your account email address or displayed in the Service. Email notices are deemed received 24 hours after sending. In-Service notices are deemed received upon first display.

Notices from you to OmniParse: Send to damr58h@gmail.com with a clear subject line (e.g., "DSR", "DPA", "Refund Request", "Account Termination"). Legal notices are deemed received on the date we reply acknowledging receipt.

22. Acceptable Use Policy

You agree not to:

- Upload documents that you do not have the right or lawful authority to process, including where doing so would violate copyright, privacy, confidentiality, contractual or other legal restrictions.

- Upload documents containing personal data of EU data subjects where you do not have an applicable legal basis for the relevant processing or, where applicable, a lawful transfer mechanism for transfers to third-country AI providers.
- Use automated scraping, bots, or rate-limit bypass tools to overwhelm the Service.
- Attempt IDOR (insecure direct object reference), prompt injection, SQL injection, XSS, or other attacks or exploits against the Service.
- Upload child sexual abuse material (CSAM), content that unlawfully promotes violence or terrorism, or other content whose processing or transmission is prohibited by applicable law in your jurisdiction or the operator's jurisdiction (Czech Republic).
- Use the Service for fraud, money laundering, tax evasion, or any other unlawful activity.

23. Contact

For questions about these Terms: **damr58h@gmail.com**

Compliant with EU Directive 93/13/EEC (unfair contract terms), EU Directive 2011/83/EU (consumer rights, including right of withdrawal), EU Directive (EU) 2023/2673 (one-click withdrawal button, effective 19 June 2026), Czech Act No. 89/2012 Coll. (Civil Code, including §§ 419, 1752, 1753, 1829, 1837, 2913), Czech Act No. 634/1992 Coll. (Consumer Protection, including §§ 16, 20b–20e), Czech Act No. 110/2019 Coll. (Czech GDPR implementation), Czech Act No. 127/2005 Coll. § 89 (Electronic Communications — ePrivacy implementation), and Czech Act No. 480/2004 Coll. (Information Society Services — unsolicited communications).

Privacy Policy

Last updated: September 19, 2026

1. Data Controller

The data controller for personal data processed through this Service is **Simon Curda** (natural person), based in the Czech Republic. For GDPR purposes (EU Regulation 2016/679), the Controller is responsible for ensuring compliance with data protection law. The Controller has not appointed a Data Protection Officer (DPO) as none is legally required under GDPR Art. 37 for an operation of this scope and nature; all GDPR-related matters are handled directly by the Controller.

Contact: **damr58h@gmail.com** (use this for all GDPR-related requests: access, rectification, erasure, restriction, portability, objection, and consent withdrawal). We aim to respond within 30 days as required by GDPR Art. 12(3). Where a request is complex or numerous, this period may be extended by a further two months pursuant to GDPR Art. 12(3), in which case we will inform you of the extension and the reasons within the initial 30-day period.

Identity of the Controller and contact details (GDPR Art. 13(1)(a)–(b)): Simon Curda, natural person, domicile Czech Republic. Email: damr58h@gmail.com. For Data Subject Requests, please use the subject line "DSR — [your name]". For B2B / DPA matters, please use the subject line "DPA — [company name]".

2. Personal Data We Collect

Account information: When you register, we collect your name and email address. Your password is hashed using bcrypt (12 rounds) and never stored in plaintext.

Documents you upload: When you use the document parsing feature, your files (PDF, JPEG, PNG, WebP) are sent to AI processing providers (see Section 6 for the full list). The original file is stored in your account database (base64-encoded) for 30 days so you can preview it in the dashboard, then deleted from the active environment after 30 days. The extracted data is stored in your account until you delete it or terminate your account.

Chat history: Messages you send to the AI chat assistant and the responses returned are stored in your account database until you delete them via the dashboard's Clear button or by deleting your account.

Technical data: We collect standard server logs including your IP address, browser type, and request timestamps for security and service stability purposes.

Local storage data: The application uses your browser's local storage to remember your authentication token, theme preference, and cookie consent choice. Invoice and chat data is stored server-side in your account, not in local storage.

3. Legal Basis for Processing (GDPR Art. 6)

- **Account creation and authentication:** Contractual necessity (Art. 6(1)(b)) — necessary to provide the Service you requested.
- **Document parsing:** Consent (Art. 6(1)(a)) — you actively upload documents for processing. You can withdraw this consent by deleting your invoices.
- **Chat interactions:** Consent (Art. 6(1)(a)) — you actively send messages to the AI chat.
- **Terms acceptance:** Consent (Art. 6(1)(a)) — recorded at signup with timestamps.
- **Security, service stability, and abuse prevention:** Legitimate interest (Art. 6(1)(f)) — server logs for abuse prevention, rate limiting, fraud detection, and account monitoring. We monitor account activity patterns (invoice volume, chat usage, email scanning frequency, account age) to detect potential abuse. Accounts suspected of abuse may be frozen (AI features blocked) or deleted in accordance with our Terms of Service. A legitimate-interest balancing test has been conducted and is documented in our internal Records of Processing Activities (ROPA it is available to supervisory authorities on request).
- **Local storage preferences:** Consent via cookie banner (Art. 6(1)(a) + ePrivacy Directive Art. 5(3), implemented in Czech Act No. 127/2005 Coll. § 89).
- **Processing of personal data appearing in customer-uploaded invoices** (e.g., names of vendor employees, contact persons, signatories): Legitimate interest (Art. 6(1)(f)) — necessary to provide the document-extraction Service the customer requested. Where such data reveals special categories (Art. 9), processing is limited to what is strictly necessary for invoice extraction and is not used for any other purpose.

4. Data Retention (GDPR Art. 5(1)(e))

We retain personal data only for as long as necessary for the purposes described in this Policy, or as required by applicable law. The following retention periods apply:

Data category Retention period Basis / exception **Uploaded document files** (PDFs, images, base64-encoded in primary DB) 30 days in active environment Auto-purged after 30 days. Extraction results remain in account. **Extracted invoice data** Until account deletion (or user deletes individual invoice) Deleted within 90 days of account termination. **Chat history** Until cleared by user or account deletion Cascade-deleted with account. **Audit logs** (invoice edits, approvals, deletions) Lifetime of account Cascade-deleted with account. **Account data** (email, name, hashed password, settings) Until account deletion User can delete at any time from Settings. **Server logs** (Vercel — IP, browser, timestamps) 30 days Vercel Hobby/Pro plans auto-delete after 30 days. **Database backups** (Supabase PITR) 7 days (rolling) Disaster recovery only; overwritten on rolling basis. **Stripe billing records** Per Stripe policy (~10 years) Required by financial/tax regulations. Governed by Stripe's privacy policy. **Incident logs** (security incidents involving AI systems) 2 years Best practice for regulatory audit purposes (AI Act Art. 12). **Withdrawal / consent records** (termsAcceptedAt, withdrawalAcknowledgedAt, withdrawnAt, lastTosEmailSentAt) Lifetime of account + 90 days Evidence of consent for regulatory audit; cascade-deleted with account.

No indefinite retention. We do not retain personal data indefinitely. Where longer retention is required by law (e.g., financial records under Czech accounting regulations), retention is limited to the statutory period and data is then deleted.

Relationship between 30-day and 90-day periods: Uploaded document *files* (the binary PDF/image) are purged after 30 days. *Extracted data* (the structured JSON — vendor, invoice number, line items) remains in your account for as long as your account is active, and is deleted within 90 days of account termination. These are different data categories with different retention needs; the 30-day period applies to file binaries, the 90-day period applies to post-termination cleanup of all remaining personal data.

5. Your Rights Under GDPR (Art. 12-23)

You have the following rights regarding your personal data:

- **Right of access (Art. 15):** Request a copy of all personal data we hold about you. Use the "Export my data" feature in Settings, or contact us.
- **Right to rectification (Art. 16):** Request correction of inaccurate personal data. Edit your profile from Settings.
- **Right to erasure / "right to be forgotten" (Art. 17):** Request deletion of your personal data. Use the "Delete account" feature in Settings. The deletion request immediately initiates the account-deletion process; active application data is deleted according to the applicable deletion schedule, while backups and legally retained records may persist for limited periods as described in §4.
- **Right to restriction of processing (Art. 18):** Request that we limit how we use your data. Contact us to request this.
- **Right to data portability (Art. 20):** Receive your data in a structured, machine-readable format. Use the "Export my data" feature in Settings.
- **Right to object (Art. 21):** Object to processing based on legitimate interests (e.g. server log collection).
- **Right to withdraw consent (Art. 7(3)):** Where OmniParse relies on consent for a specific processing activity, you may withdraw that consent at any time using the mechanism provided for that activity or by contacting us. Withdrawal does not affect the lawfulness of processing carried out before withdrawal. Withdrawal of consent does not by itself require deletion of data where another lawful basis for continued processing applies. The right to erasure is described separately above.
- **Rights related to automated decision-making and profiling (Art. 22):** OmniParse does **not** engage in solely automated decision-making producing legal or similarly significant effects concerning you. AI

extraction and chat responses are tools to assist you; all material decisions (e.g., invoice approval, payment authorization) require your explicit human action. You have the right not to be subject to a decision based solely on automated processing; if you believe any such decision has occurred, contact us immediately.

To exercise any of these rights, contact us at **damr58h@gmail.com**. We will respond within 30 days as required by GDPR Art. 12(3). You also have the right to lodge a complaint with a supervisory authority — in particular the Czech Office for Personal Data Protection (ÚOOÚ) at **uouu.cz** (<https://www.uouu.cz>), or, if you reside in another EU/EEA Member State, the supervisory authority of your habitual residence.

6. Subprocessors and AI Providers

We use third-party service providers to provide the Service. Depending on the processing activity and provider, a provider may act as a processor/subprocessor on our behalf or as an independent controller for its own legally defined purposes. Where a provider processes personal data on our behalf, we maintain appropriate contractual data-protection arrangements as required by applicable law. Where personal data is transferred to a third country, we use an applicable GDPR Chapter V transfer mechanism where required.

The current provider configuration and applicable processing locations are described below.

- **Vercel Inc.** (United States, DPF-certified) — Web hosting and serverless function execution. Processes: IP addresses, request metadata. **Vercel DPA** (<https://vercel.com/legal/dpa>).
- **Supabase Inc.** (Ireland, EU) — Primary database hosting (PostgreSQL). Processes: all stored account data, invoices, chat messages. DPA built into Terms of Service (effective August 1, 2026) — automatically applies to all customers, no separate signature required. **Supabase DPA** (<https://supabase.com/legal/customer-resources/data-processing-addendum>). Subprocessor change notifications: **supabase.com/legal/customer-resources/subprocessor-list** (<https://supabase.com/legal/customer-resources/subprocessor-list>).
- **Mistral AI** (Paris, France — EU) — Primary AI inference provider. Processes: uploaded document content (transiently) for vision extraction, chat messages (transiently). **EU-based — no SCC required**. Training opt-out (`usage_options.enable_training=false`) is sent on every request when `MISTRAL_DISABLE_TRAINING=true` (default). **Mistral Privacy Policy** (<https://mistral.ai/legal/privacy-policy>).
- **Groq Inc.** (United States) — Secondary AI inference provider (after Mistral). Processes: uploaded document content (transiently), chat messages (transiently). **SCCs confirmed** (EU SCC Module 2, self-executing upon acceptance of Groq Services Agreement; DPA dated October 15, 2025). DPA accessible at **console.groq.com** (<https://console.groq.com>). Governing law: Ireland. *Note: Groq is optional in our configuration — operators may deploy in EU-only mode (Mistral only) by leaving GROQ_API_KEY unset.*
- **Google LLC** (United States, DPF-certified) — Final fallback AI inference provider. Models used: Gemini 2.0 Flash, Gemini 2.5 Flash, Gemini 1.5 Flash. **DISABLED BY DEFAULT** — uses Google's AI Studio free-tier endpoint (`generative_language.googleapis.com`), whose data-handling terms are weaker than Google Cloud Vertex AI (which has a self-executing DPA). Operator must set `ENABLE_GOOGLE_GEMINI=true` after reviewing Google AI Studio terms. For production with EU personal data, operators are strongly encouraged to migrate to Vertex AI instead. Google LLC itself is DPF-certified. **Google Cloud DPA** (<https://cloud.google.com/terms/data-processing-addendum>) (applies to Vertex AI path, not AI Studio free tier).
- **Stripe Inc.** (United States, DPF-certified) — Payment processing for paid plans. Processes: email, billing details. Stripe may process payment and transaction information for its own legally defined purposes,

including payment security, fraud prevention and regulatory compliance. Depending on the processing activity, Stripe may act as an independent controller rather than solely as a processor on behalf of OmniParse. Stripe's own privacy and cookie notices apply to processing carried out by Stripe. **Stripe DPA** (<https://stripe.com/legal/dpa>).

OpenRouter — DISABLED by default. OpenRouter is an AI inference aggregator that routes requests to underlying model providers (Meta, Nvidia, Google, etc.). OpenRouter's free-tier models (all model IDs ending in :free) typically permit training on prompt content, and there is no DPA or SCC in place with OpenRouter as of September 2026. **OpenRouter is therefore disabled by default in our production configuration** (the `ENABLE_OPENROUTER` environment variable must be explicitly set to `true` by the operator after completing their own DPA/SCC review). When disabled, OpenRouter is never called — document content and chat messages are never transmitted to OpenRouter or its underlying providers.

AI vision cascade (production default): Mistral (EU, primary) → Groq (US, SCCs confirmed, secondary). OpenRouter and Google Gemini are **disabled by default** and only used if the operator explicitly enables them after completing their own DPA/SCC review (OpenRouter) or AI Studio terms review (Google Gemini).

SCC Status (as of September 18, 2026): **Supabase Inc.** (Ireland, EU) — DPA built into Terms of Service (effective August 1, 2026). Automatically applies to all customers. EU-based — no SCC required. **Vercel Inc.** (United States) — DPF-certified. DPA available at vercel.com/legal/dpa. **Groq Inc.** — Groq is used as a secondary AI provider in the default production cascade. OmniParse relies on the applicable Groq contractual safeguards for international transfers, including SCCs where applicable (EU SCC Module 2, self-executing upon acceptance of Groq Services Agreement; DPA dated October 15, 2025). Groq's operational logging, retention and security practices may include limited processing for reliability, security or abuse prevention. OmniParse does not represent that provider-side retention is zero in all circumstances. The applicable provider terms and configuration govern. DPA accessible at console.groq.com (<https://console.groq.com>). Governing law: Ireland. Competent authority: Irish DPC. 72-hour breach notification. DPF certification status pending verification. **Mistral AI** — Established in the European Union (Paris, France). OmniParse will verify the applicable processing locations, subprocessors and remote-access arrangements for the specific Mistral service used. Where a restricted transfer outside the EEA occurs, OmniParse will implement an appropriate GDPR Chapter V transfer mechanism and supplementary safeguards where required. **Google LLC (Gemini)** — Google Gemini through the Google AI Studio free-tier endpoint is **disabled by default** and is not used for customer content in the default production configuration. The Google Cloud DPA and Vertex AI contractual terms do not automatically apply to the Google AI Studio free-tier endpoint. OmniParse will not enable this provider for customer content until the applicable service terms, data-use conditions, processing locations and transfer safeguards have been reviewed and documented. If Google-based processing is required for EU personal data, OmniParse will evaluate an appropriate enterprise service and contractual configuration, such as Vertex AI, before enabling it. Migration does not by itself establish compliance; the actual configuration and applicable agreements must be verified. **Stripe Inc.** (United States) — DPF-certified. DPA available at stripe.com/legal/dpa. **OpenRouter — DISABLED by default.** Not used in production unless operator explicitly sets `ENABLE_OPENROUTER=true` after completing DPA/SCC review. When disabled, no data is transferred to OpenRouter or its underlying providers.

EU users: AI processing via Mistral (EU-based) and Groq (US-based, SCCs confirmed) has appropriate safeguards in place under GDPR. Google is DPF-certified. OpenRouter is disabled by default. However, users should note that documents uploaded to the Service may contain personal data of multiple data subjects (e.g. vendor names, employee names, email addresses, bank details). The user uploading such documents is responsible for ensuring they have a valid legal basis under GDPR Art. 6 for processing that data, and that the

transfer to AI providers is lawful under Chapter V.

Model training — production configuration: OmniParse sends `usage_options={"enable_training": false}` on applicable Mistral API requests. This parameter does not, by itself, guarantee that all data-use conditions are identical across Mistral service tiers. OmniParse will use a Mistral configuration for customer content only after verifying the applicable contractual terms, training controls and processing conditions. Google Gemini (AI Studio free tier, whose data-handling terms are weaker than Google Cloud Vertex AI) is disabled by default. OpenRouter (which has no training opt-out on free-tier models) is also disabled by default. Groq's data-handling terms are governed by the Groq DPA and SCCs. Operators who change this configuration must update this Policy accordingly.

Important — transfer mechanism availability: The availability of transfer mechanisms under GDPR Chapter V may change as a result of regulatory guidance, court decisions or changes in provider arrangements. If a transfer mechanism relied upon by OmniParse becomes unavailable or insufficient, OmniParse will reassess the affected transfer and implement an appropriate lawful mechanism and supplementary safeguards where required before continuing the affected processing. The current status of court proceedings, certifications and contracts is maintained as internal compliance evidence with verification dates, and is reflected in this Policy as of the "Last updated" date above.

7. Security Measures (GDPR Art. 32)

OmniParse implements appropriate technical and organizational measures to ensure a level of security appropriate to the risk, taking into account the state of the art and the costs of implementation. Measures include:

- **Encryption in transit:** HTTPS/TLS 1.2+ for all connections (Vercel-managed certificates, auto-renewed).
- **Encryption at rest:** Supabase (PostgreSQL) provides transparent disk encryption (TDE) on managed infrastructure. IMAP credentials stored by users are encrypted with AES-256-GCM.
- **Password storage:** Passwords are hashed with bcrypt (12 rounds) and never stored in plaintext or reversible form.
- **Authentication:** JWT-based authentication with 7-day token expiry. Tokens are signed with a server-side secret and rotated periodically.
- **Rate limiting:** 5 login attempts per minute per IP on authentication endpoints; broader rate limits on AI endpoints to prevent abuse.
- **Security headers:** Content-Security-Policy, Strict-Transport-Security (HSTS, 2-year max-age, preload), X-Frame-Options: DENY, X-Content-Type-Options: nosniff, Referrer-Policy: strict-origin-when-cross-origin.
- **Per-user data isolation:** Every database query includes a `userId` filter — no insecure direct object reference (IDOR) is possible at the query layer.
- **Prompt-injection defense:** All chat messages are screened against 30+ prompt-injection attack patterns before being sent to AI providers.
- **Account deletion:** Transactional cascade delete ensures full erasure of all user-owned rows, including invoices, chat, audit logs, and IMAP credentials.
- **Vendor due diligence:** Subprocessors are assessed before engagement and re-assessed at least annually. DPAs and (where applicable) SCCs are in place (see §6 and §11).

- **Incident response:** A documented incident-response procedure exists. Where a personal-data breach affecting customer data occurs, OmniParse will notify the affected customer as described in §11.9 and reasonably cooperate with the customer in fulfilling applicable obligations under GDPR Articles 33 and 34. The customer (as controller) is responsible for notifying the competent supervisory authority (in the Czech Republic, the ÚOOÚ) and affected data subjects where required.

8. Children's Data (GDPR Art. 8)

The Service is intended for business and professional users and is not directed at children. Users must be at least **18 years old** to register an account, unless applicable law permits use by a younger person with the required parental or guardian consent and the contractual arrangement is legally valid. In the Czech Republic, the age of digital consent under GDPR Art. 8 (as implemented by Act No. 110/2019 Coll.) is 15; in other EU/EEA Member States the age ranges from 13 to 16 depending on national law.

Where applicable law permits use by a person aged between the applicable digital-consent age and 18, that person must obtain parental or guardian consent to use the Service and to enter into these Terms. Where a minor accesses the Service through an organizational account (e.g., an employee or contractor of a business customer), the organization is responsible for ensuring that such access is lawful and complies with applicable child-labor and data-protection law.

If we learn that we have collected personal data from a child below the applicable age of digital consent without verifiable parental consent, we will delete that data promptly. Contact us at damr58h@gmail.com if you believe a child has provided us with personal data.

9. Czech Republic Compliance

For users in the Czech Republic, the following applies in addition to GDPR:

- **Act No. 110/2019 Coll.** on the processing of personal data — Czech implementation of GDPR. Sets the age of digital consent at 15 (Art. 8 GDPR) and codifies national derogations.
- **Act No. 89/2012 Coll.** (Civil Code) — contractual relationships, consumer protection, liability regime (incl. § 2913 on limitation of liability).
- **Act No. 127/2005 Coll.** (Electronic Communications Act, "Lex Telecom") § 89 — implements the ePrivacy Directive (2002/58/EC) cookie/storage consent rule, as amended effective 1 January 2022. Requires opt-in consent before storing or accessing information on a user's terminal equipment, except where strictly necessary (Art. 5(3) exception).
- **Act No. 480/2004 Coll.** (Information Society Services Act) — governs unsolicited commercial communications (spam) and certain information-society-service provider obligations. *Note: this Act is not the cookie-consent statute; that role belongs to Act No. 127/2005 Coll. § 89.*
- **Act No. 634/1992 Coll.** (Consumer Protection Act) §§ 16, 20b–20e — out-of-court dispute resolution (ADR) before Česká obchodní inspekce (ČOI).
- **Digital Services Act (Regulation (EU) 2022/2065):** Directly applicable in the Czech Republic since 17 February 2024. Depending on the features made available and the manner in which users interact with and share content, particular aspects of the Service may fall within relevant provisions of the Digital Services Act. OmniParse does not provide a public social-networking or marketplace functionality in its intended default use. The applicability of specific DSA obligations depends on the actual service design and applicable legal classification.

- **Supervisory authority (data protection):** Úřad pro ochranu osobních údajů (Office for Personal Data Protection), Pplk. Sochora 27, 170 00 Praha 7, Czech Republic. Website: [uoou.cz](https://www.uoou.cz) (<https://www.uoou.cz>).
- **Consumer-protection authority:** Česká obchodní inspekce (ČOI), Štěpánská 567, 120 00 Praha 2. Website: coi.gov.cz (<https://coi.gov.cz>).

10. Canadian PIPEDA Compliance

For users accessing the Service from Canada, OmniParse complies with the Personal Information Protection and Electronic Documents Act (PIPEDA, R.S.C. 2000, c. P-8.9). Bill C-27 (which would have replaced PIPEDA with the Consumer Privacy Protection Act) died on the Order Paper on 6 January 2025 and did not become law; PIPEDA therefore remains in force unamended.

The ten fair-information principles in Schedule 1 of PIPEDA are respected, in particular:

- **Accountability:** Simon Curda (natural person, Czech Republic) is responsible for personal information under PIPEDA and has designated damr58h@gmail.com as the contact for privacy matters.
- **Identifying purposes:** Purposes are described in §2 of this Policy.
- **Consent:** Where consent is required under PIPEDA, OmniParse obtains consent through the applicable signup, service or privacy mechanism. Acceptance of the Terms of Service is not, by itself, treated as consent for every separate processing purpose. Consent for non-essential browser storage is obtained through the cookie/storage consent mechanism.
- **Limiting collection:** We collect only what is necessary for the Service (see §2).
- **Limiting use, disclosure, retention:** Retention periods in §4 apply.
- **Accuracy:** Users can correct their data via Settings; AI-extracted data can be edited in the dashboard.
- **Safeguards:** See §7 (security measures).
- **Openness:** This Policy and the Cookie Policy are publicly available without account.
- **Individual access:** Canadian users may request access to, correction of, or deletion of their personal data by emailing damr58h@gmail.com.
- **Challenging compliance:** Canadian users may challenge our compliance by contacting us or by filing a complaint with the Office of the Privacy Commissioner of Canada at [priv.gc.ca](https://www.priv.gc.ca) (<https://www.priv.gc.ca>).

Mandatory breach notification (PIPEDA § 10.1 and Breach of Security Safeguards Regulations): Where a breach of security safeguards involving personal information under our control creates a real risk of significant harm to a Canadian individual, we will notify the affected individual and the Privacy Commissioner of Canada as soon as feasible after determining the breach has occurred, and we will keep a record of the breach for at least 24 months.

11. Data Processing Addendum (DPA)

This Data Processing Addendum ("DPA") applies when you (the "Customer") act as a data controller and use the Service to process personal data on behalf of data subjects. This DPA reflects the parties' agreement with respect to the processing of personal data under GDPR Art. 28.

11.1 Roles, Scope, and Confidentiality

- **Customer** is the data controller. Customer determines the purposes and means of processing.
- **OmniParse** (operated by Simon Curda) is the data processor. We process personal data only on documented instructions from Customer.

- We process documents and extracted data per your instructions (via the Service's UI and API). We do not use your data for any purpose other than providing the Service.
- **Confidentiality:** OmniParse will ensure that persons authorized to process personal data are subject to an obligation of confidentiality under contract or applicable law. Such persons will process personal data only as necessary to provide, secure and support the Service or as otherwise required by applicable law.

11.2 Processor Obligations (GDPR Art. 28(3))

Pursuant to GDPR Art. 28(3), OmniParse will:

- (a) process personal data only on documented instructions from the Customer, including with regard to transfers of personal data to a third country or an international organisation, unless processing is required by applicable Union or Member State law;
- (b) ensure that persons authorised to process personal data are subject to an appropriate obligation of confidentiality under contract or applicable law;
- (c) implement appropriate technical and organisational measures to ensure a level of security appropriate to the risk, as described in §11.5;
- (d) comply with the conditions applicable to engaging subprocessors, including the authorisation and notification mechanisms described in §11.3 and §11.8;
- (e) taking into account the nature of the processing, reasonably assist the Customer through appropriate technical and organisational measures in responding to data-subject rights requests, as described in §11.4;
- (f) taking into account the nature of the processing and the information available to OmniParse, reasonably assist the Customer with obligations under GDPR Articles 32 to 36, including security measures, personal-data breach notification, data-protection impact assessments and prior consultation where applicable;
- (g) make available to the Customer information reasonably necessary to demonstrate compliance with the obligations applicable to processors under GDPR Art. 28 and allow for and contribute to audits and inspections as described in §11.10;
- (h) at the Customer's choice, delete or return personal data after the end of the provision of the relevant processing services, unless applicable law requires continued storage, as described in §11.7.

11.3 Subprocessors

OmniParse engages the following subprocessors. Customer grants general written authorization for OmniParse to engage these subprocessors; the current list is maintained in this Privacy Policy §6 and material changes will be notified 30 days in advance per §11.8:

Active subprocessors in the default production configuration:

- **Supabase Inc.** (Ireland, EU) — Primary database hosting (PostgreSQL)
- **Vercel Inc.** (United States, DPF-certified) — Web hosting and serverless function execution
- **Mistral AI** (Paris, France — EU) — Primary AI processing
- **Groq Inc.** (United States, SCCs confirmed) — Secondary AI processing
- **Stripe Inc.** (United States, DPF-certified) — Payment processing, where applicable

Optional providers disabled by default:

- **Google Gemini** (AI Studio free tier) and **OpenRouter** are not used for customer content in the default production configuration. They may only be enabled following an internal review of applicable terms,

processing locations, transfer mechanisms and data-use conditions. The subprocessor list and relevant customer notices will be updated before such processing is activated, where required.

Note: If a provider is enabled without updating the contractual terms or notices applicable to an existing B2B customer, the customer's authorization and notification rights under §11.8 must be addressed. Mere toggling of an environment variable is not, by itself, legal authorization.

11.4 Data Subject Rights

Customer is responsible for responding to data subject rights requests (access, rectification, erasure, portability, objection). OmniParse will assist Customer with such requests, including by exporting or deleting personal data upon Customer's written request. Contact **damr58h@gmail.com** with the subject "DPA — Data Subject Request".

11.5 Security Measures (GDPR Art. 32)

OmniParse implements the following technical and organisational measures, subject to the Service configuration and applicable infrastructure-provider controls:

- Encryption in transit: HTTPS/TLS 1.2+ for connections to the Service and applicable third-party providers.
- Encryption at rest: Supabase-managed encryption for database storage. IMAP credentials stored by OmniParse are encrypted using AES-256-GCM before storage.
- Password security: User passwords are hashed using bcrypt and are not stored in plaintext.
- Authentication: JWT-based authentication with a configured token-expiry period.
- Per-user data isolation: Application queries are scoped to the authenticated user's account, with user-level access controls applied at the application layer.
- File retention: Uploaded file binaries are automatically purged from active application storage according to the retention schedule in the Privacy Policy.
- Account deletion: Account deletion triggers deletion of user-owned application data according to the applicable deletion and backup schedule.
- Audit logging: Significant account and invoice actions are logged for security and operational purposes.
- Rate limiting: Authentication and AI-related endpoints are subject to rate limits designed to reduce abuse.
- Security headers: Security headers including CSP, HSTS, X-Frame-Options, X-Content-Type-Options and Referrer-Policy are configured where applicable.
- Prompt-injection screening: Chat messages are screened against known prompt-injection patterns before being sent to AI providers.
- Vendor due diligence: Subprocessors are assessed before engagement and periodically thereafter, taking into account the nature and risk of the processing.

11.6 International Transfers (GDPR Chapter V)

Personal data may be transferred to the United States for processing by Groq and, only if explicitly enabled, by other US-based providers.

Where a US recipient is covered by an applicable adequacy mechanism, OmniParse may rely on that mechanism. Where an adequacy mechanism does not apply, OmniParse will rely on an appropriate GDPR Chapter V safeguard, such as Standard Contractual Clauses, together with supplementary measures and a Transfer Impact Assessment where appropriate.

Provider status as of September 19, 2026:

- **Mistral AI** — EU-based, Paris, France. Processing that remains within the EEA does not constitute a third-country transfer under GDPR Chapter V.
- **Groq Inc.** — US-based. OmniParse relies on the contractual and transfer safeguards applicable under the Groq services agreement and DPA, including EU SCCs where applicable.
- **Google LLC** — disabled by default. Google AI Studio free-tier processing is governed by Google's applicable AI Studio terms. The Google Cloud Data Processing Addendum applies to the applicable Google Cloud / Vertex AI services and is not treated as automatically covering Google AI Studio free-tier processing.
- **Vercel Inc.** — US-based. The applicable contractual and transfer safeguards are described in Vercel's data-processing documentation.
- **Stripe Inc.** — US-based. Stripe's applicable contractual, transfer and controller/processor arrangements depend on the specific payment-processing activity.
- **OpenRouter** — disabled by default. No customer content is transferred to OpenRouter unless the operator explicitly enables it after completing the applicable contractual, data-protection and transfer review.

The cascade prioritizes Mistral first, then Groq — both have the applicable safeguards described above. Google Gemini and OpenRouter are disabled by default and are not used for customer content in the default production configuration.

11.7 Duration and Deletion

This DPA continues for the duration of the Customer's use of the Service. Following termination of the processing services, OmniParse will, at the Customer's documented request and where technically available, return or make available for export personal data processed on the Customer's behalf, or delete such personal data in accordance with the retention and deletion schedule below, unless applicable law requires continued retention.

File binaries are deleted from active application storage within 30 days. Other personal data associated with the account is deleted within 90 days of account termination, except where longer retention is required by law or where limited retention is necessary for security, legal or backup purposes. Backup copies are removed according to the applicable backup-rotation schedule.

Any personal data retained after termination remains subject to applicable confidentiality, security and purpose-limitation obligations and will not be processed for unrelated purposes.

11.8 Subprocessor Authorization

Customer grants general written authorization for OmniParse to engage the subprocessors listed in Section 11.3. OmniParse will notify Customer of any planned changes to the subprocessor list at least 30 days in advance via the Service or by email. Customer may object to a new subprocessor in writing within 30 days of notification. In the event of a reasonable objection, OmniParse will use commercially reasonable efforts to make available an alternative or recommend a change to Customer's use of the Service to avoid processing by the objected-to subprocessor. If OmniParse is unable to provide a reasonable alternative, Customer may suspend or terminate the affected Service.

11.9 Personal Data Breach Notification (GDPR Art. 33-34)

OmniParse will notify the Customer without undue delay after becoming aware of a personal-data breach affecting personal data processed on the Customer's behalf. The notification will include, to the extent available, the nature of the incident, affected data and systems, likely consequences, mitigation measures and a contact point for further information. OmniParse will reasonably cooperate with the Customer in fulfilling applicable obligations under GDPR Articles 33 and 34, taking into account the nature of processing and available information.

11.10 Audit Rights

Audits will normally be conducted no more than once per calendar year, upon reasonable prior notice and during normal business hours. This limitation will not apply where an audit is reasonably necessary following a personal-data breach, a material compliance concern, a binding request from a competent supervisory authority, or where required by applicable law. The parties will seek to minimize disruption and protect confidential information. Alternatively, OmniParse may provide a current third-party audit report (e.g., SOC 2 Type II, ISO 27001) or equivalent documentation in lieu of an on-site audit, which the parties agree is a reasonable alternative under GDPR Art. 28(3)(h), where OmniParse has such a report available.

11.11 Acceptance and Activation

For B2B customers (businesses, organizations, accounting firms): This DPA takes effect and is binding between OmniParse (as processor) and the B2B Customer (as controller) upon the later of: (a) the Customer's acceptance of the Terms of Service, and (b) the Customer's first use of the Service to process personal data. The DPA is incorporated by reference into the Terms of Service; no separate signature is required for it to take effect (per GDPR Art. 28(9) and Recital 81, which permit DPAs to be formed by written agreement including electronic form).

Activation mechanism: B2B Customers who require a standalone signed DPA (e.g., for their own compliance records, customer-facing DPAs, or large Enterprise contracts) may request one by contacting **damr58h@gmail.com** with the subject line "DPA Activation — [Company Name]". We will execute a standalone Data Processing Agreement within 10 business days. Until then, this embedded DPA (Privacy Policy §11) governs the processing relationship.

For individual users: Where a user processes personal data in a professional or organizational capacity, the user acts as the controller and OmniParse acts as the processor where OmniParse processes that data on the user's documented instructions. Where processing is exclusively for personal or household activities, the applicability of the GDPR household exemption will depend on the circumstances. The parties' actual roles are determined by the nature and purposes of the processing and not solely by the wording of this DPA.

Version control: The "Last updated" date at the top of this Privacy Policy reflects the current DPA version. Material changes to this DPA will be notified to B2B Customers at least 30 days in advance per §11.8 (Subprocessor Authorization mechanism). Prior versions are retained and available on request.

12. Privacy Notice for Invoice Data Subjects (GDPR Art. 14)

Where customers upload documents (e.g., invoices) that contain personal data of individuals other than themselves — for example, contact persons at vendor companies, employee names, signatories, or bank-account holders — that personal data is processed by OmniParse on the customer's instructions. The customer is the controller with respect to that data; OmniParse acts as a processor on the customer's behalf.

Pursuant to GDPR Art. 14, and in reliance on the Art. 14(5)(b) exemption (providing the information would involve a disproportionate effort because OmniParse has no direct relationship with those data subjects and they may be numerous), this public Privacy Notice serves as the transparency mechanism. We ask our customers to inform their data subjects of the use of OmniParse where feasible and to direct them to this Notice.

If your personal data appears in a document processed by OmniParse:

- **Who is processing your data:** The OmniParse customer who uploaded the document is the data controller. OmniParse (Simon Curda, natural person, Czech Republic) acts as a processor on their behalf.
- **Purposes:** Extraction of structured data from invoices for accounting, accounts-payable, and vendor-management purposes.
- **Legal basis:** The applicable legal basis is determined by the customer as data controller. Legitimate interests under GDPR Art. 6(1)(f) may be applicable for ordinary business invoice-processing activities, but the customer is responsible for determining and documenting the appropriate legal basis for its specific processing.
- **Recipients:** Depending on the customer's configuration and the Service's applicable production configuration, personal data may be processed by AI inference providers and infrastructure providers listed in §6. Providers disabled by default are not used for customer content unless the applicable configuration is explicitly changed.
- **Retention:** Retention depends on the type of data and the applicable retention schedule in Section 4. When a customer deletes a document, OmniParse will delete or anonymize associated data within the applicable operational deletion period, subject to legally required retention, security records and backup-cycle limitations. The customer account termination schedule is described separately in Section 11.7.
- **Your rights:** You have rights of access, rectification, erasure, restriction, portability, and objection. To exercise these rights, contact the customer (the controller) directly; if you cannot reach them, contact us at **damr58h@gmail.com** and we will route your request.
- **Right to complain:** You may lodge a complaint with the Czech DPA (ÚOOÚ) or with the supervisory authority of your habitual residence.

13. Changes to This Policy

We may update this Privacy Policy from time to time. Material changes will be notified via the Service or by email at least 30 days before taking effect. The "Last updated" date at the top of this page indicates when this Policy was last revised. Continued use of the Service after changes constitutes acceptance of the updated Policy.

14. Contact

- **Data controller:** Simon Curda, natural person, Czech Republic
- **General contact & all GDPR requests:** damr58h@gmail.com
- **B2B / DPA matters:** damr58h@gmail.com (subject line "DPA — [company name]")
- **Data Subject Requests (Art. 15-22):** damr58h@gmail.com (subject line "DSR — [your name]")
- **Czech supervisory authority:** uoou.cz (<https://www.uoou.cz>)
- **Canadian Commissioner:** priv.gc.ca (<https://www.priv.gc.ca>)

Cookie Policy

Last updated: September 19, 2026

1. What Are Cookies and Local Storage

Cookies are small text files stored on your device by websites you visit. OmniParse primarily uses browser Local Storage rather than traditional cookies. Under the EU ePrivacy Directive (2002/58/EC) as amended by Directive 2009/136/EC, the consent requirements apply to both cookies and Local Storage. Each item stored is classified as either essential or non-essential based on its purpose.

2. What We Store and Why

Essential storage (always active — exempt from consent per ePrivacy Art. 5(3), as implemented by Czech Act No. 127/2005 Coll. § 89):

- **Authentication token (op_token):** Stores your JWT session token in Local Storage so you stay signed in across page refreshes. *Retention:* 7 days (matching JWT expiry), then auto-cleared on next page load. *Purpose:* Strictly necessary for the Service to function.
- **Cookie consent record (omniparse_cookie_consent):** The consent-preference record is used to remember the user's storage preferences and to avoid repeatedly displaying the consent prompt. Its classification and retention are assessed according to its specific purpose and applicable law. *Retention:* 12 months from last update, then re-prompted. *Purpose:* Prevents re-showing the consent banner and records your consent decision for accountability.

Non-essential storage (requires consent — only set after you accept via the cookie banner):

- **Theme preference (theme):** Stores your light/dark mode choice. *Retention:* Until withdrawn or browser data cleared. *Purpose:* User preference (UI customization).
- **Keyboard shortcuts (op_shortcuts):** Stores custom keyboard shortcut assignments. *Retention:* Until withdrawn or browser data cleared. *Purpose:* User preference (UI customization).
- **Legal consent (op_legal_consent):** Records GDPR data-transfer consent for document uploads. *Retention:* Duration of account; refreshed on each upload if revoked. *Purpose:* Compliance (GDPR consent evidence). *Note on separation:* This storage item records a consent decision; it is distinct from the legal basis for processing documents, which is determined separately based on the user's role, the purpose of processing and applicable law. Where consent is the legal basis, it will be requested through a separate, specific and informed mechanism and will not be bundled with unrelated cookie preferences.
- **Crash logs (op_crash_log):** Stores recent client-side errors for debugging. *Retention:* Rolling window — last 10 errors only, FIFO overwrite. *Purpose:* Service stability and debugging. Client-side crash logs are limited to a rolling window of the latest 10 entries. OmniParse is designed to avoid including document contents, authentication tokens or unnecessary personal data in crash logs. Where personal data is inadvertently included, it is handled under the applicable retention and security controls.

Analytics:

- Currently, no analytics cookies or storage items are active. If we implement analytics in the future, we will (a) update this Policy with the specific item name, purpose, and retention; (b) seek your explicit consent via the cookie banner; and (c) ensure the analytics provider is GDPR-compliant (e.g., DPF-certified or SCC-backed).

Third-party cookies:

- OmniParse does not set third-party cookies on the `omniparse-ai.vercel.app` domain. No advertising, tracking, social-media, or fingerprinting cookies are used.
- **Note about Stripe:** If you upgrade to a paid plan, Stripe's checkout page (at `checkout.stripe.com`) may set its own cookies on the Stripe domain (not `omniparse-ai.vercel.app`). These are governed by **Stripe's cookie policy** (<https://stripe.com/cookies>) and are necessary for payment security and fraud prevention. OmniParse does not control Stripe's cookies.

Other storage technologies:

- OmniParse's current production configuration uses Local Storage (not browser cookies) for client-side state. If cookies are introduced in the future (e.g., for session management, CSRF protection, or framework-level features), this Policy will be updated with the specific cookie name, purpose, duration, security attributes (Secure, HttpOnly, SameSite), and whether the cookie is first-party or third-party.

3. Legal Basis

Under the EU ePrivacy Directive (2002/58/EC), as amended by Directive 2009/136/EC, and GDPR:

- **Essential storage (`op_token` and the storage item used to remember the user's consent preference):** This storage is used where strictly necessary for the Service requested by the user or to remember the user's privacy preference. It is exempt from prior consent only to the extent that the applicable ePrivacy exception for strictly necessary storage applies.
- **Non-essential storage (`theme`, `op_shortcuts`, `op_legal_consent`, `op_crash_log`):** Requires your explicit, informed, prior consent via the cookie banner (GDPR Art. 6(1)(a) + ePrivacy Art. 5(3)). Consent must be as easy to withdraw as to give. You may withdraw consent at any time by (a) clicking the "Essential only" button in the cookie banner re-openable via Settings → Privacy, or (b) clearing your browser's Local Storage.
- You have the right to be informed about what data is collected and for what purpose (this section).
- You have the right to lodge a complaint with the Czech DPA (ÚOOÚ) if you believe our cookie practices violate applicable law.

4. Cookie Banner and Granular Consent

On your first visit to the Service, a cookie banner appears at the bottom of the screen. You have the following options:

- **"Accept all":** Accepts all non-essential storage categories currently offered by the Service. Essential storage remains active where it is strictly necessary for the requested Service or to remember your privacy preference.
- **"Essential only":** Records consent for essential storage only; non-essential items are not set, and any previously set non-essential items are cleared.
- **"Manage preferences" (planned):** A granular preference-management feature is planned. Until it is implemented, the banner provides the currently available choices described below. The Service does not represent that per-category consent controls are available before that feature is released.

The banner does not block page load (essential cookies work regardless). Your choice is recorded in Local Storage under the key `omniparse_cookie_consent` along with a timestamp, and persists for 12 months. After 12 months, or if you clear your Local Storage, the banner will re-appear.

You can change your choice at any time by clicking the "Cookie settings" link in the footer, or by clearing the `omniparse_cookie_consent` key from your browser's Local Storage.

5. Managing Your Preferences

You can manage or remove Local Storage data through your browser settings:

- **Chrome:** Settings > Privacy and security > Cookies and other site data > Site data
- **Firefox:** Settings > Privacy & Security > Cookies and Site Data
- **Safari:** Preferences > Privacy > Manage Website Data

Blocking or deleting Local Storage data may affect the Service's functionality — for example, your authentication session and theme preference will be lost.

6. Czech Republic

Under Czech Act No. 127/2005 Coll. (Electronic Communications Act) § 89, as amended effective 1 January 2022 to align with the ePrivacy Directive: the processing of Local Storage data and cookies requires the user's consent unless it is strictly necessary for the operation of the service requested by the user. OmniParse complies with this requirement through the cookie consent banner (§ 4 above). Act No. 480/2004 Coll. (Information Society Services Act) applies to unsolicited commercial communications and is not relevant to OmniParse's cookie practices, as OmniParse does not engage in direct marketing.

7. Contact

For questions about this Cookie Policy: damr58h@gmail.com

Complies with EU ePrivacy Directive 2002/58/EC (as amended by Directive 2009/136/EC), GDPR Art. 6(1)(a) and Art. 7, Czech Act No. 127/2005 Coll. § 89 (Electronic Communications — ePrivacy implementation), and EDPB Guidelines 03/2022 on dark patterns in social media platform interfaces.